

# Apache Spark

## Scalable Data Processing

# What's big data?

---

- extremely large datasets that are hard to deal using relational databases
- requires different approaches: techniques, tools and architectures
- big data generates values from storage and processing of very large quantities of digital information that cannot be analyzed with traditional techniques

## What's big data, again?

- Traditionally, computation has been processor bound
- but as data increase becomes the bottleneck



# Early solution: bigger computers

---


- faster processor, more memory
- but even this couldn't keep up

## Distributed systems

The better solution, more machines

Examples:

- google
- hadoop
- tez
- spark
- flink







# Challenges with distributed systems

---

- Programming complexity
  - Keeping data and processes in sync
- Finite bandwidth
- Partial failures

# Distributed systems: evolution

## Why Apache Flink is Next-Gen?

			
<ul style="list-style-type: none"><li>• Batch</li></ul>	<ul style="list-style-type: none"><li>• Batch</li><li>• Interactive</li></ul>	<ul style="list-style-type: none"><li>• Batch</li><li>• Interactive</li><li>• Near-Real time Streaming</li><li>• Iterative processing</li></ul>	<ul style="list-style-type: none"><li>• Batch</li><li>• Interactive</li><li>• <b>Real-Time Streaming</b></li><li>• <b>Native Iterative processing</b></li></ul>
MapReduce	<b>Direct Acyclic Graphs (DAG) Dataflows</b>	RDD: <b>Resilient Distributed Datasets</b>	<b>Cyclic Dataflows</b>
1 <sup>st</sup> Generation (1G)	2 <sup>nd</sup> Generation (2G)	3 <sup>rd</sup> Generation (3G)	4 <sup>th</sup> Generation (4G)

12



# Apache spark

---

**Apache Spark is a fast, general engine for large-scale data processing on a cluster**

- Originally developed in Berkeley
- The creators founded Databricks to commercialize Spark
- Open source Apache project
  - committer for Yahoo, Databricks, Berkeley, Intel, Cloudera...
  - one of the most active and fastest-growing Apache projects



# Advantages of spark

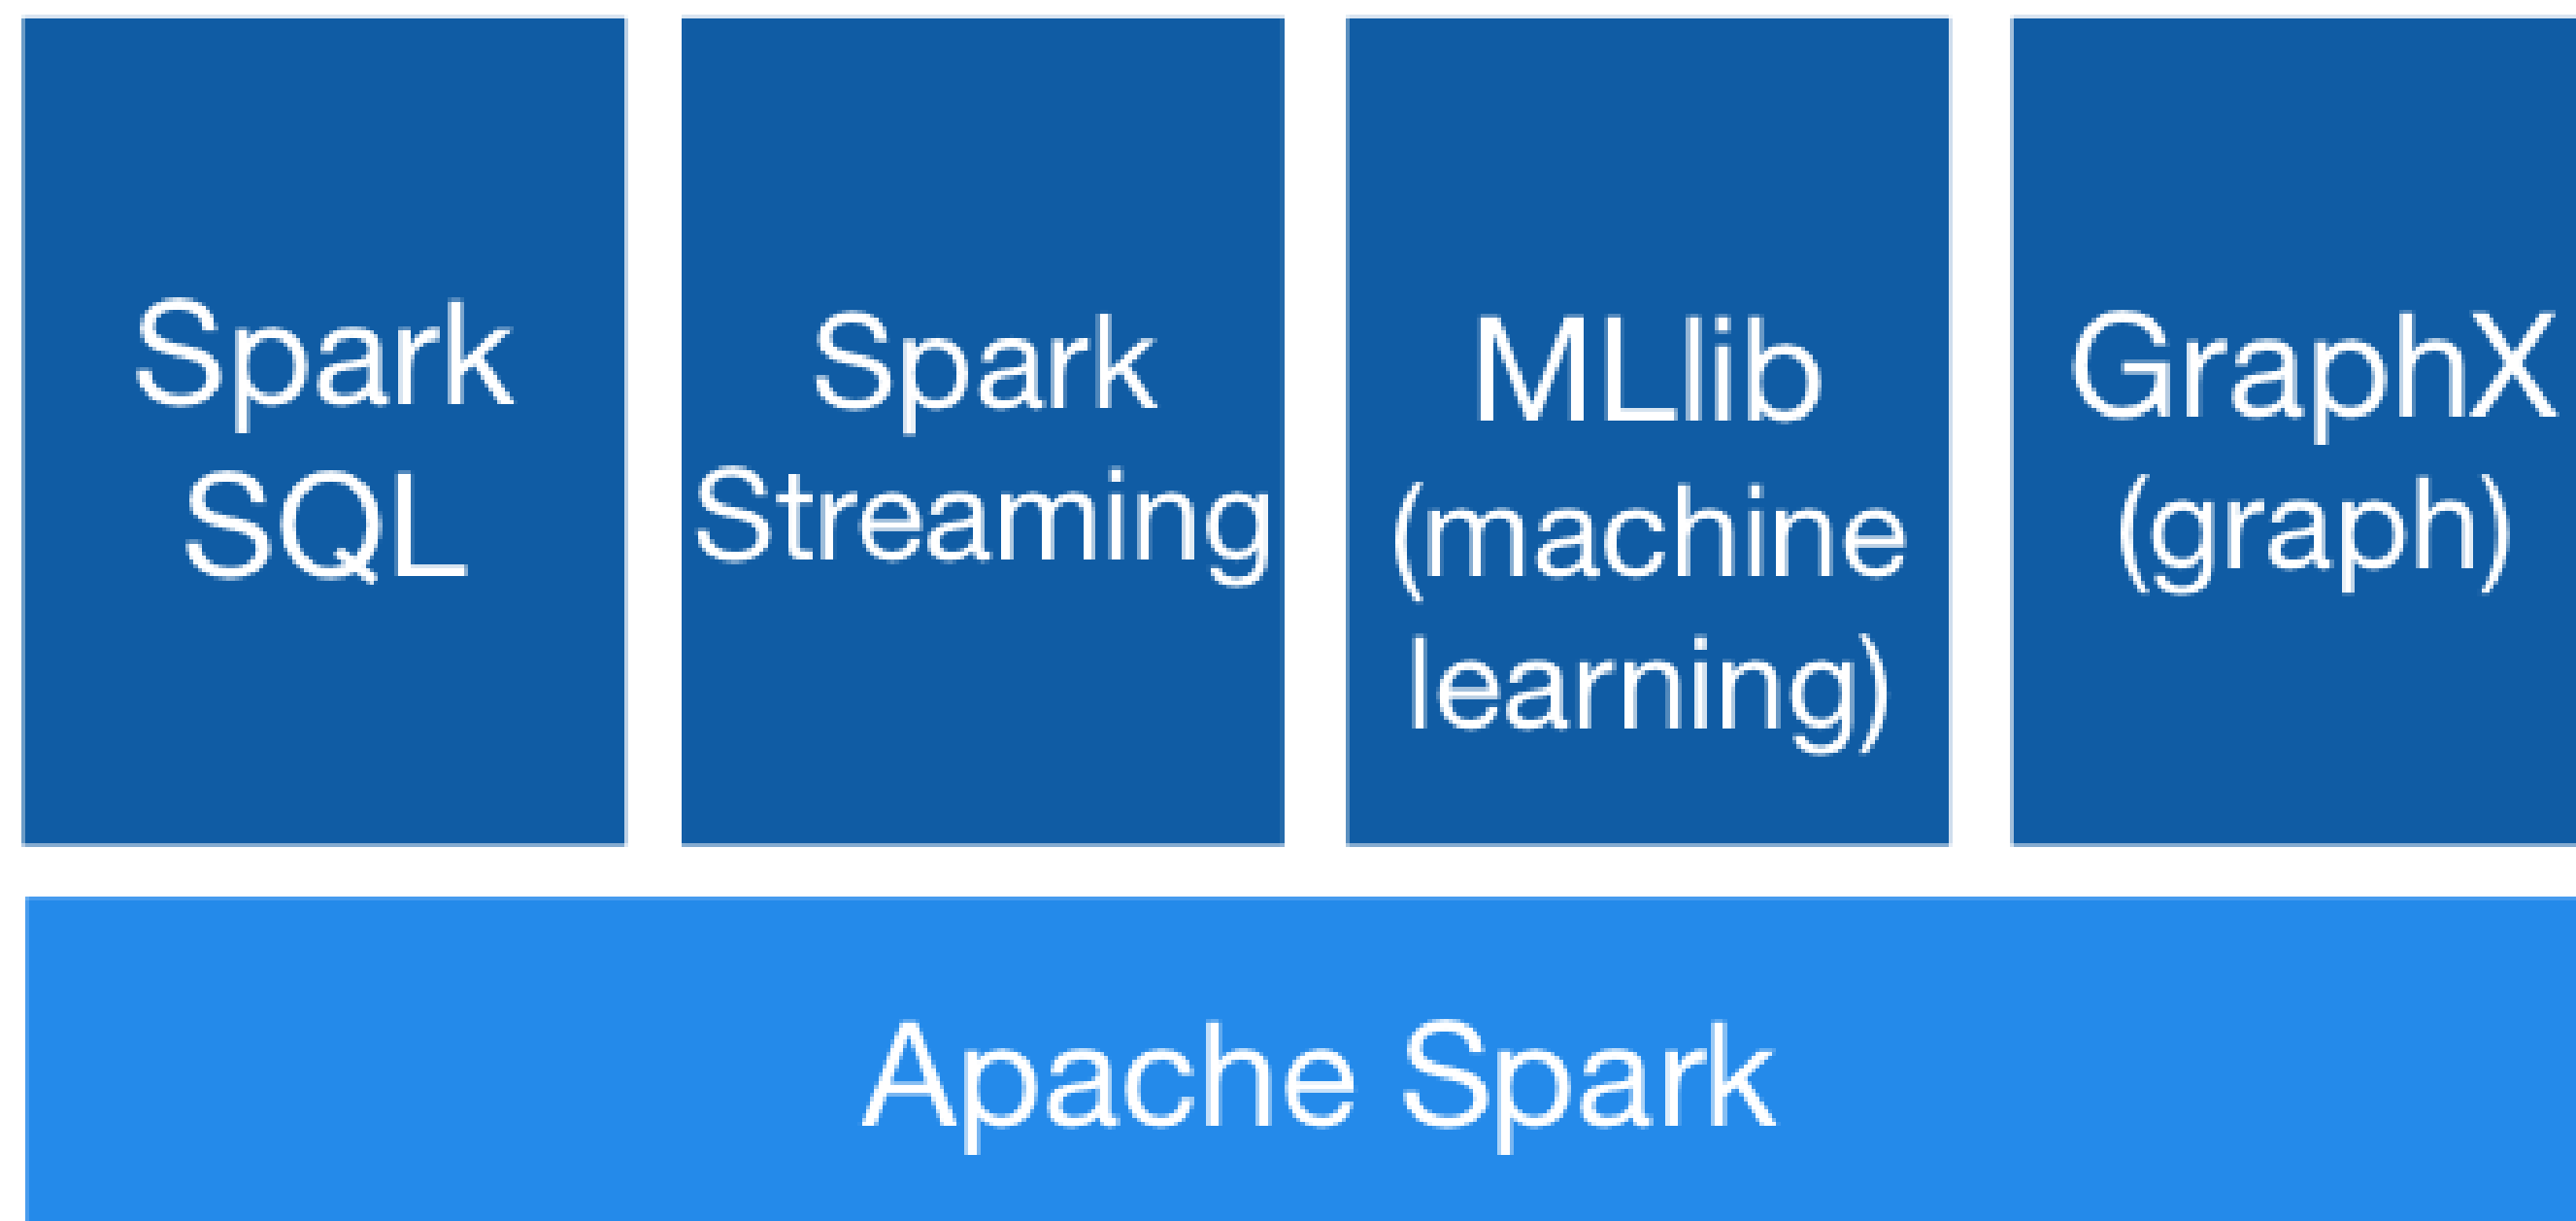
---

- High level programming framework
- Cluster computing
- Distributed storage
- Data in memory
- Provides fault tolerance
- Adding nodes adds capacity proportionally

# The spark stack

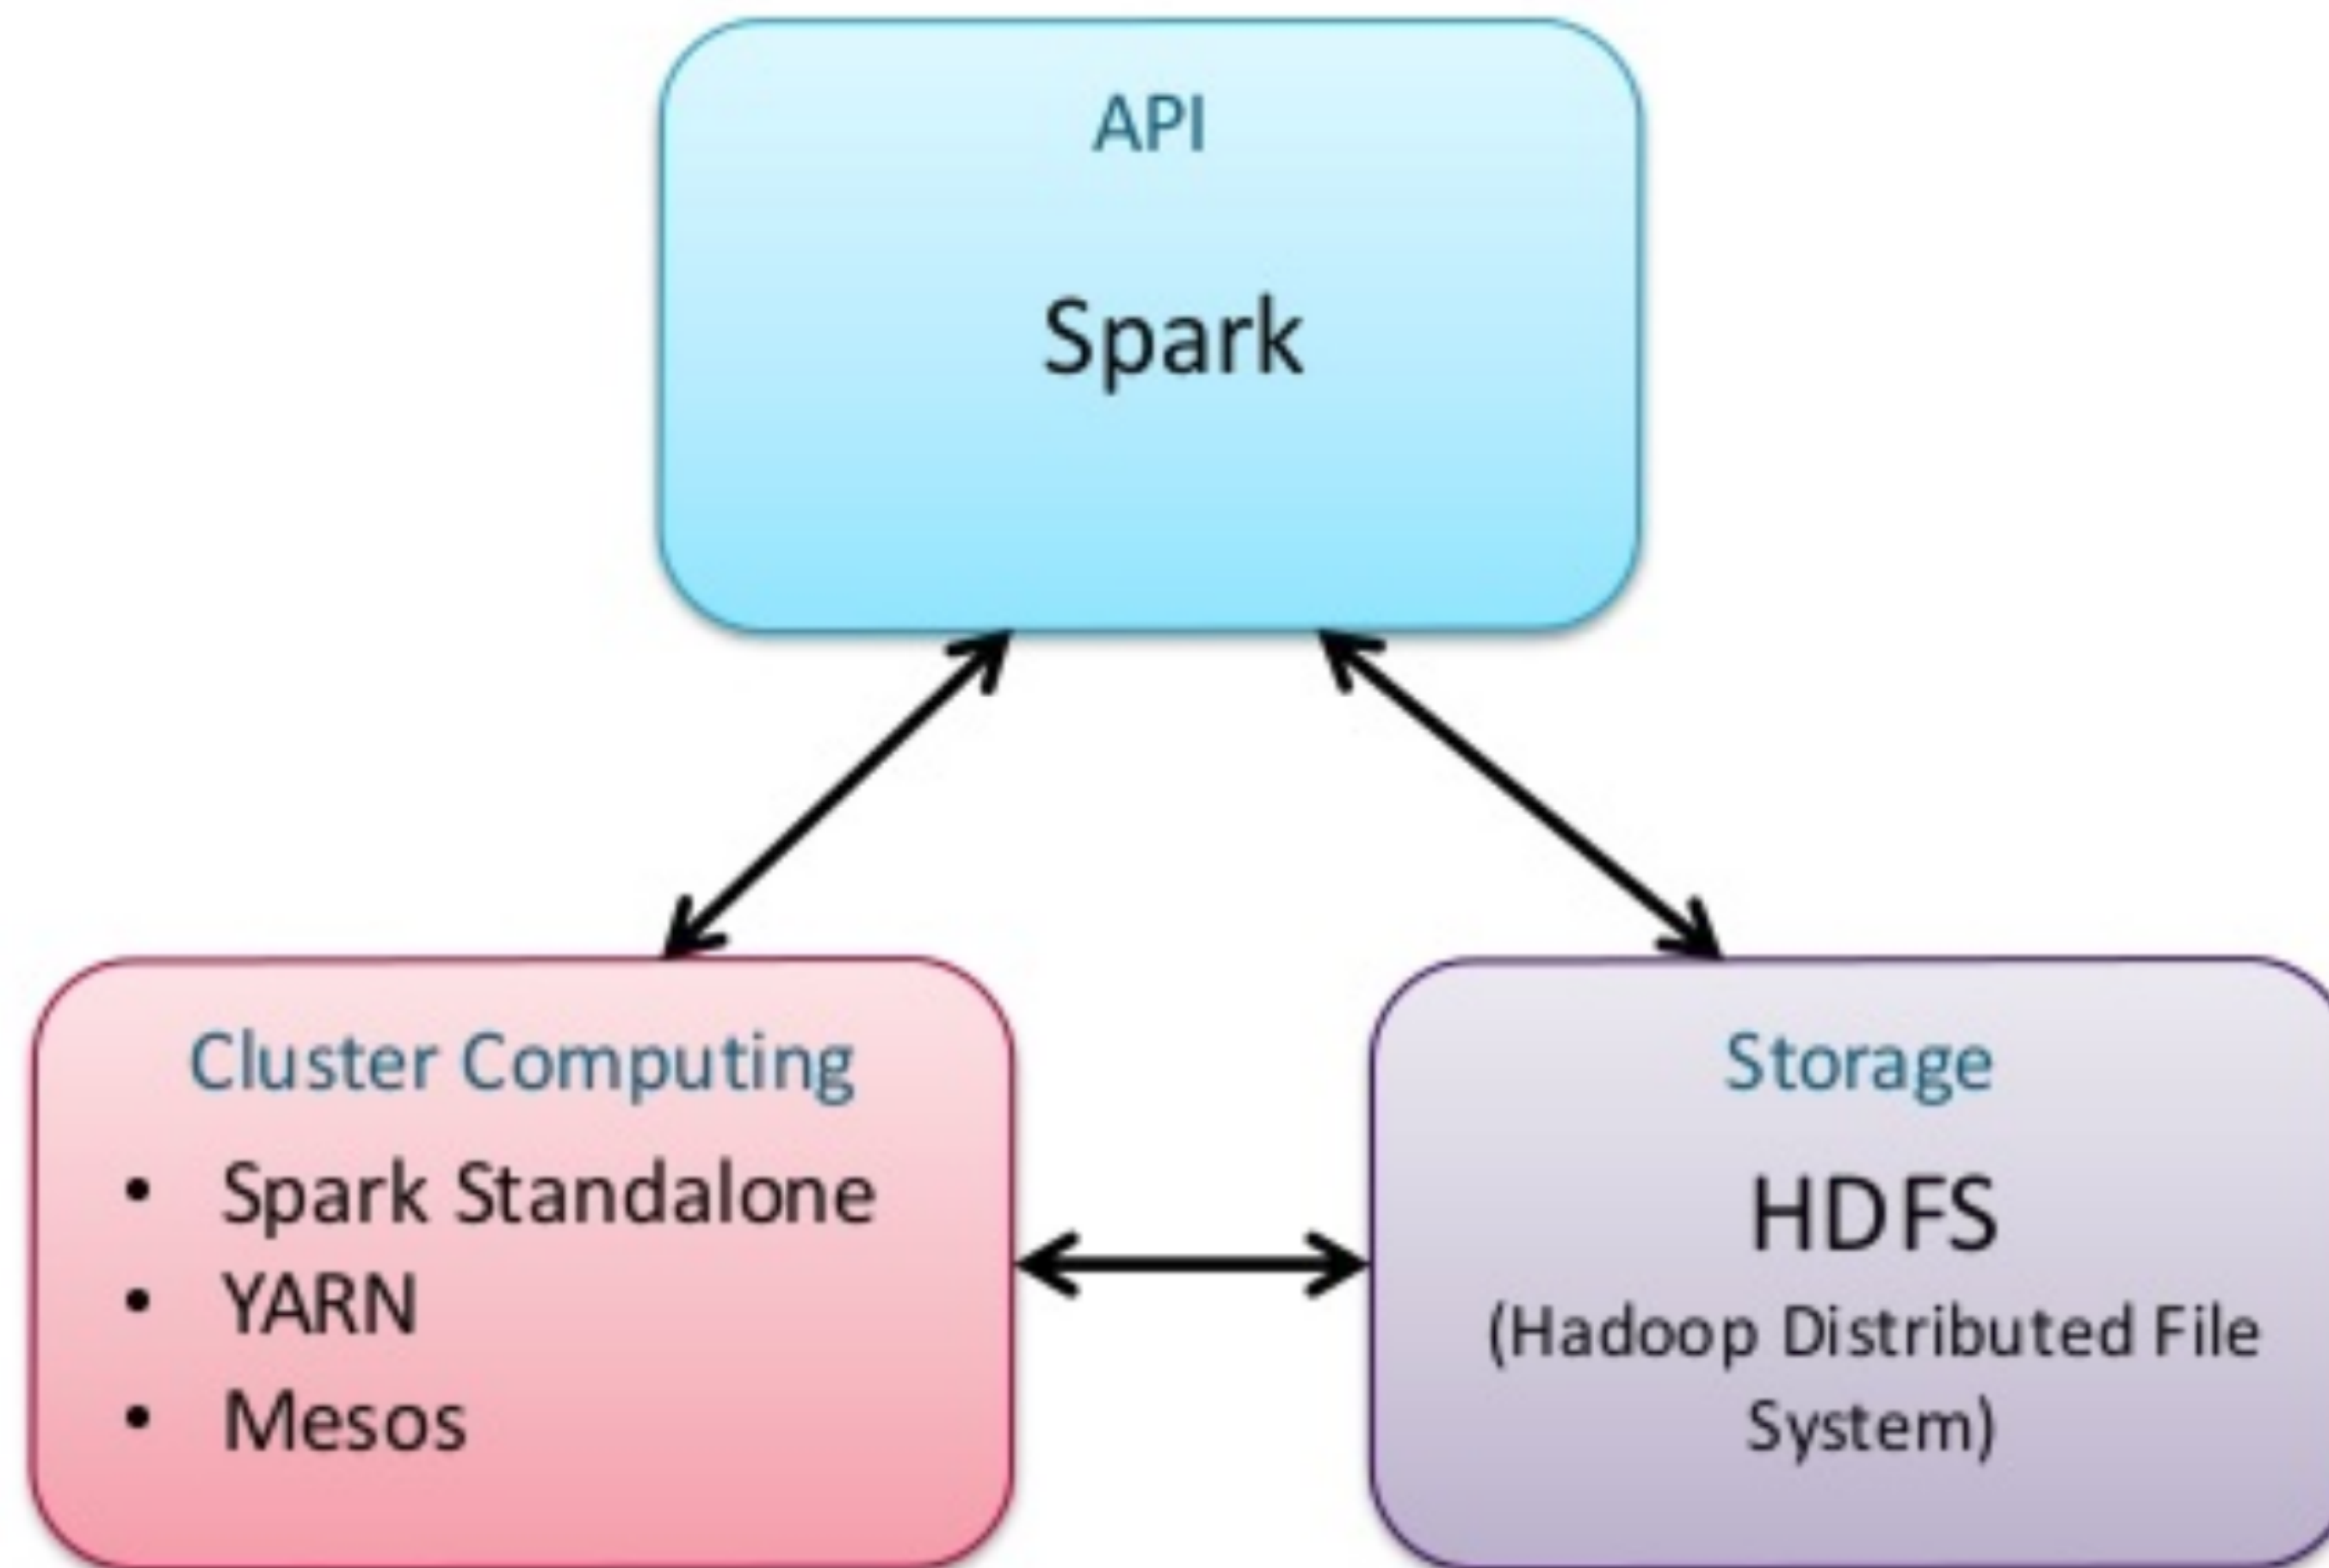
---

- Core Spark provides the fundamental Spark abstraction: Resilient Distributed Datasets (RDDs)
- Spark SQL works with structured data
- MLlib supports scalable machine learning
- Spark Streaming applications process data in real time
- GraphX works with graphs and graph-parallel computation





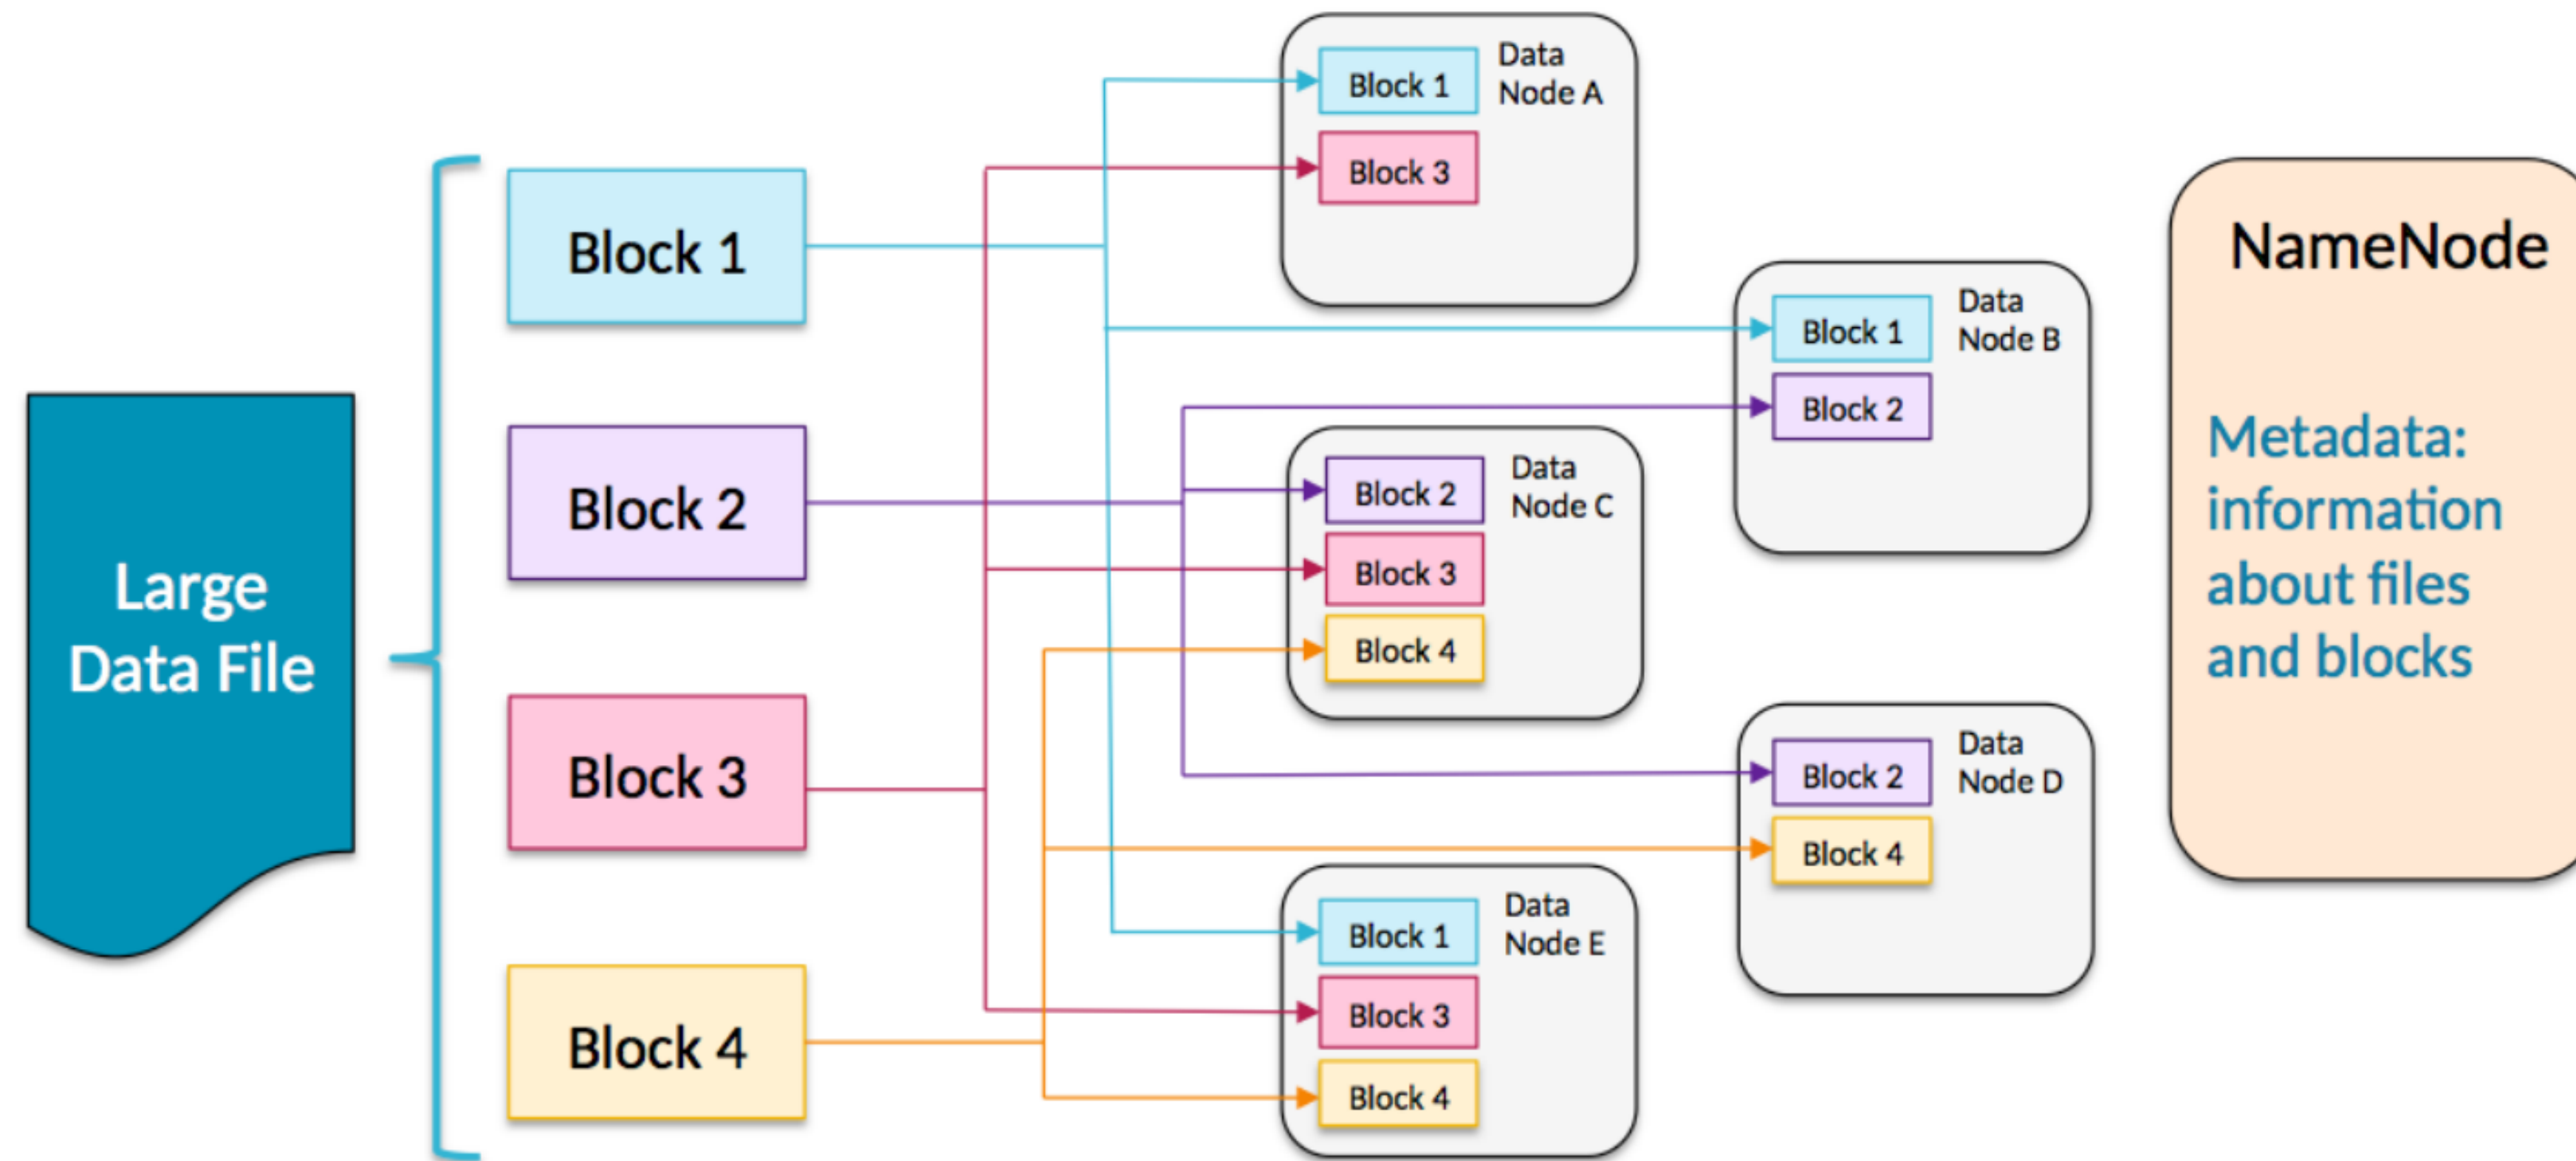
# Distributed Processing with the Spark Framework



# Architecture HDFS (1)

- HDFS is a file system written in Java. Based on Google File System
- Provides redundant storage for massive amounts of data
- Data files are split into blocks (default 128MB)
- The actual blocks are stored on cluster worker nodes (DataNodes)
- Each block is replicated
- A cluster master node runs the HDFS Name Node service (NameNode)

# Architecture HDFS (2)



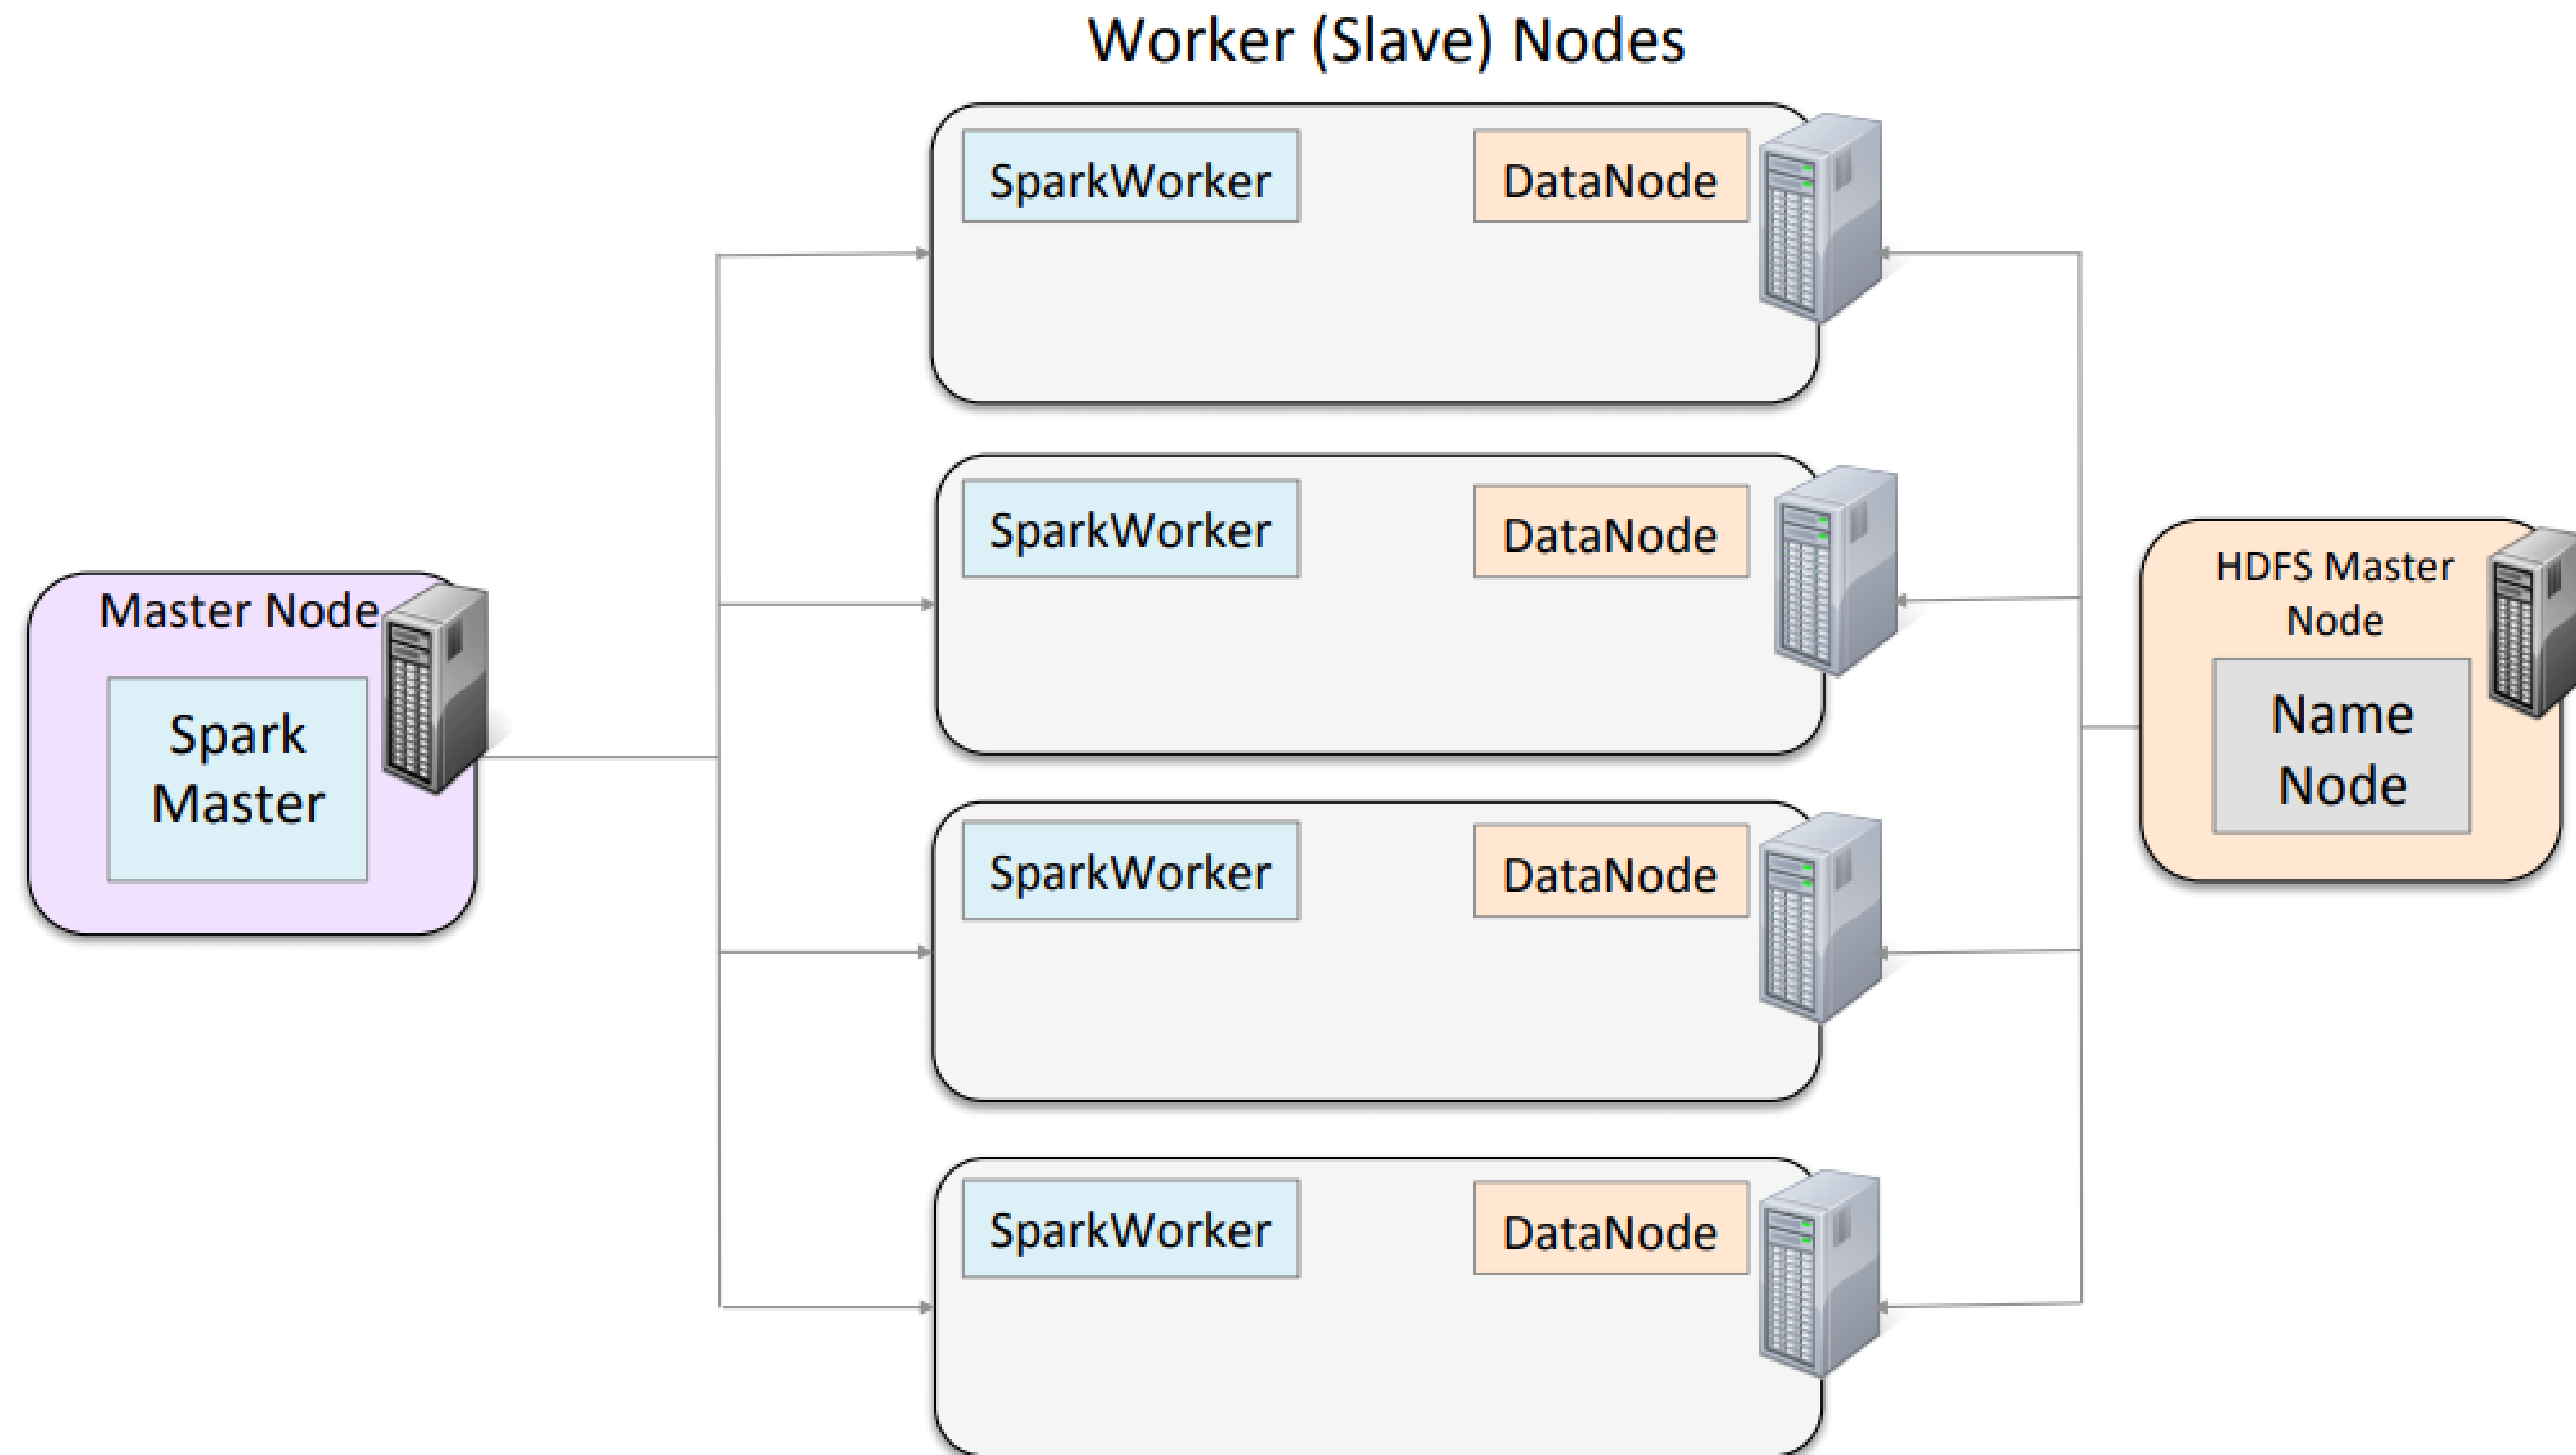
# Architecture Spark Standalone Cluster (1)

## Spark Standalone daemons

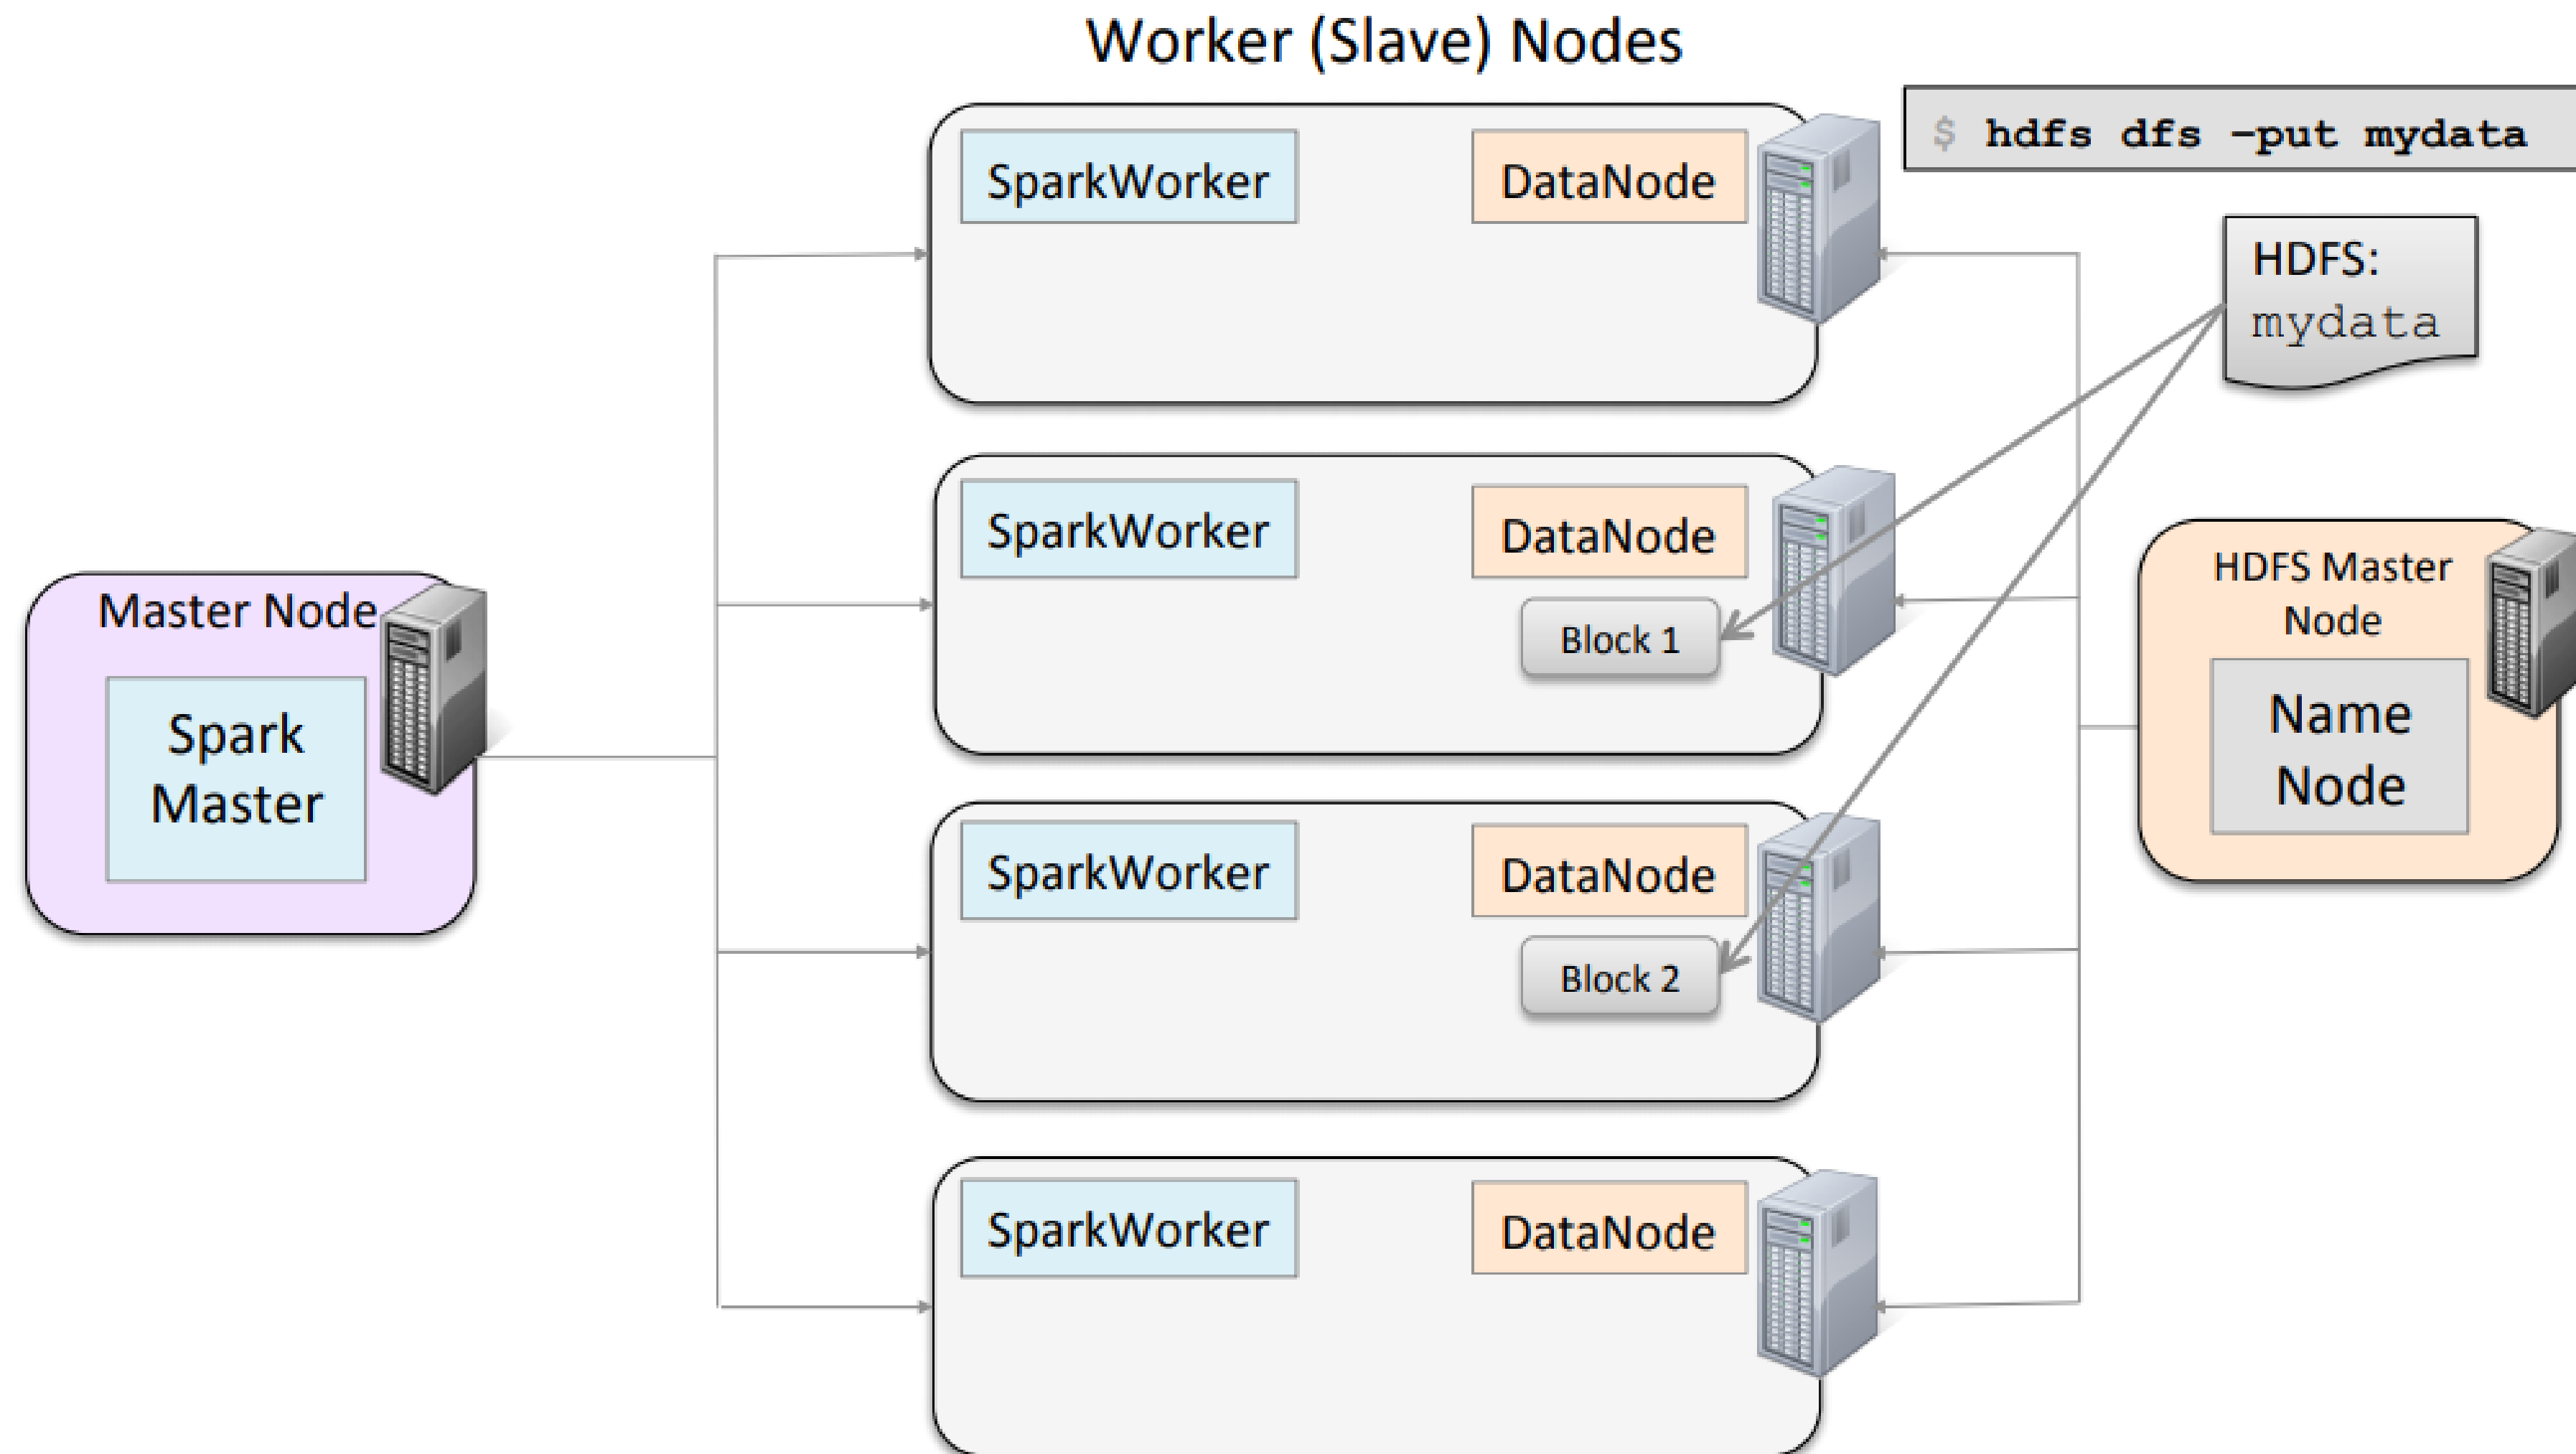
- Spark Master: One per cluster. Manages applications, distributes individual tasks to Spark Workers
- Spark Worker: One per worker node. Starts and monitors Executors for applications



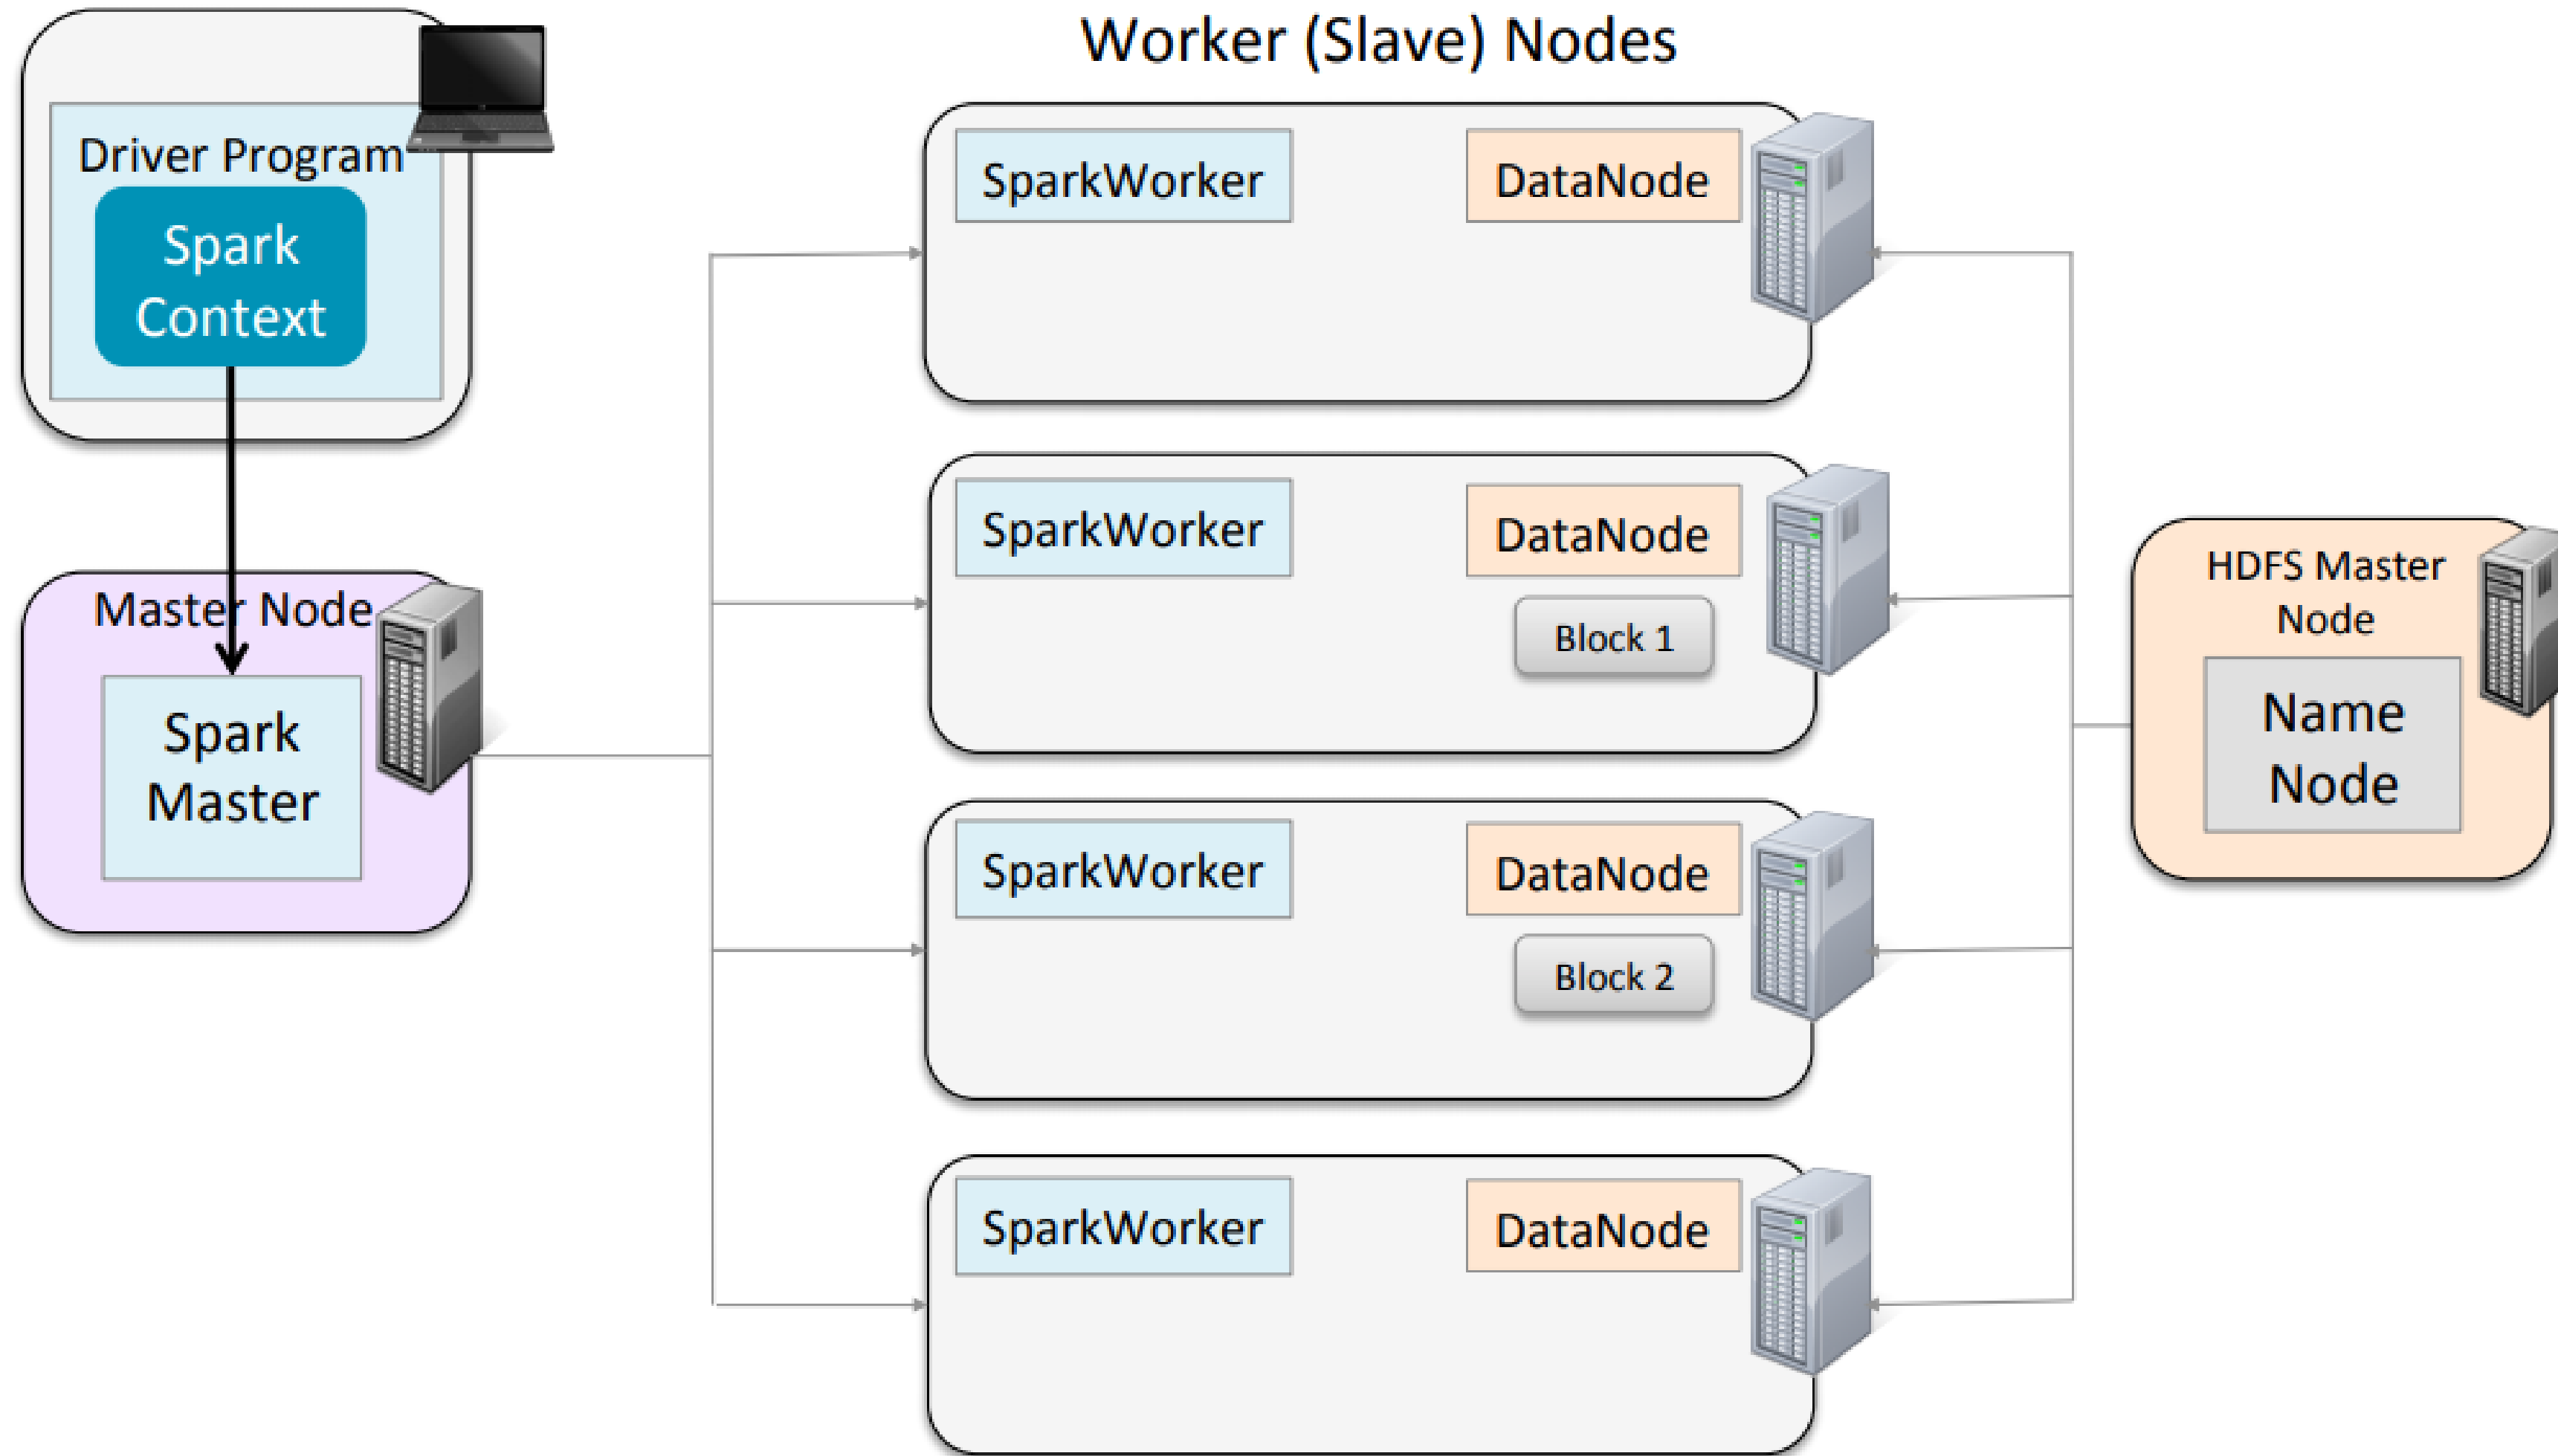
# Architecture Spark Standalone Cluster (2)



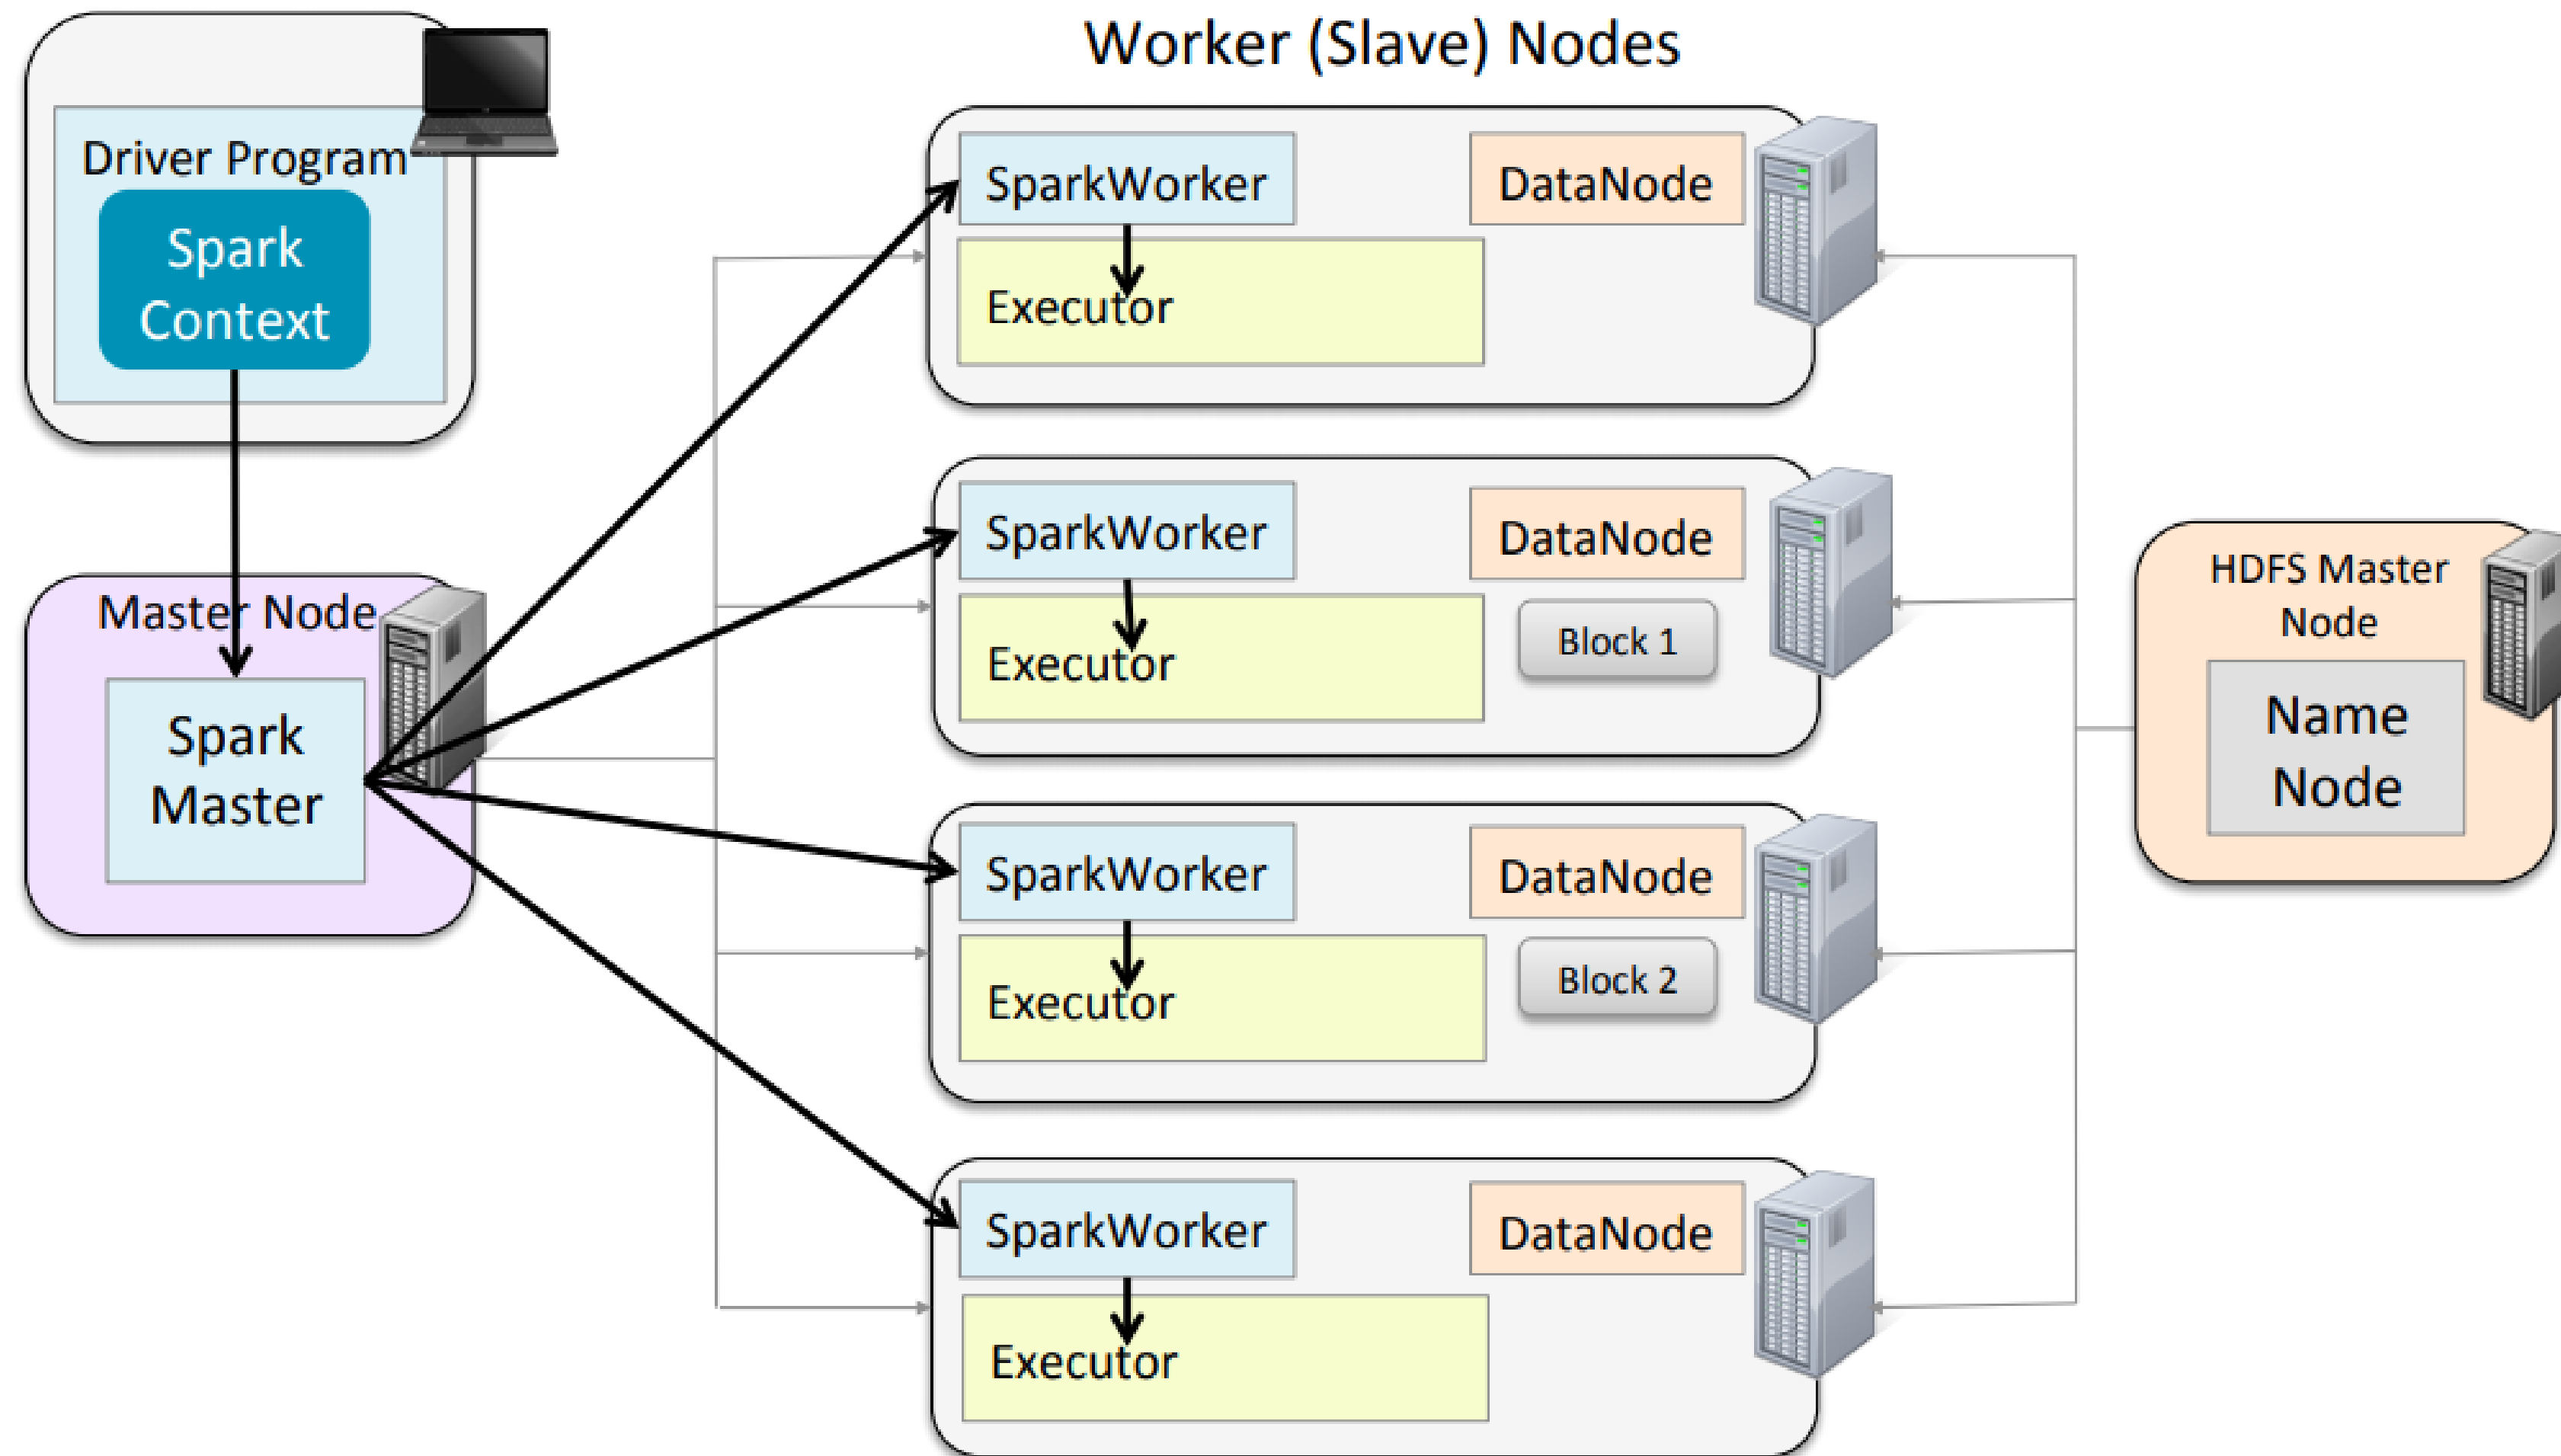
# Architecture Spark Standalone Cluster (3)



# Architecture Spark Standalone Cluster (4)

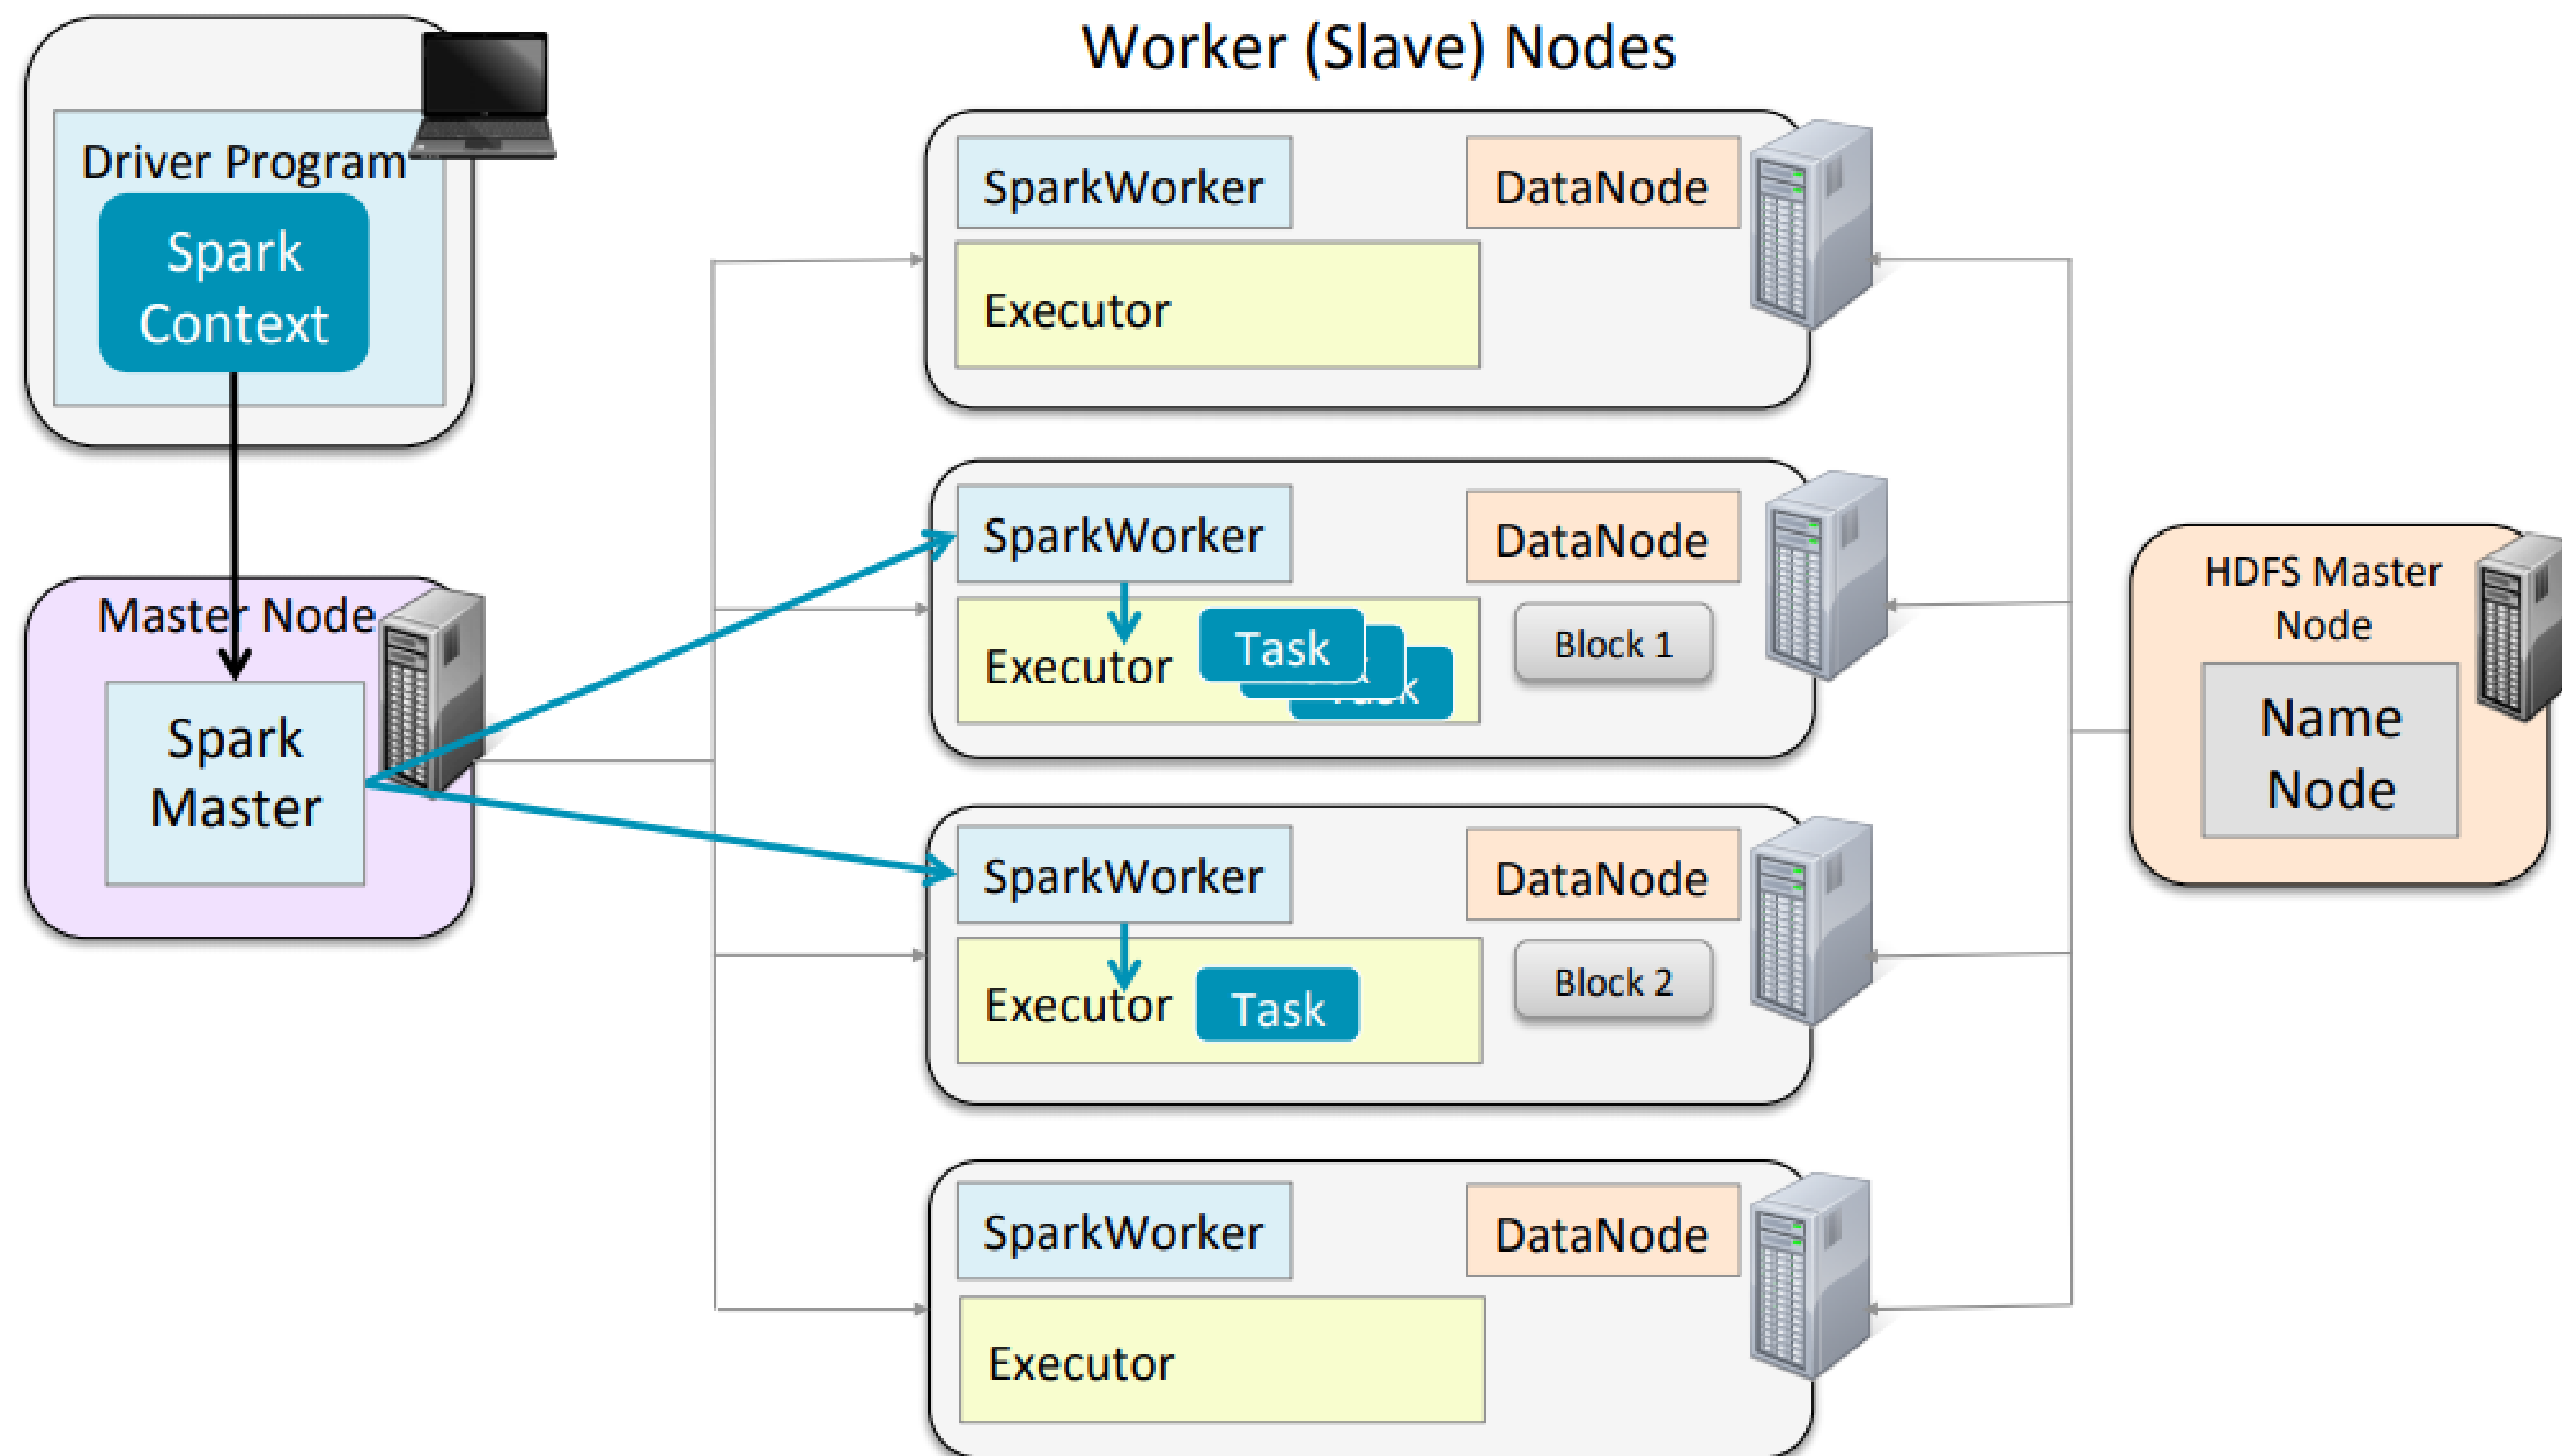


# Architecture Spark Standalone Cluster (5)





# Architecture Spark Standalone Cluster (6)



# Architecture Yarn Cluster (1)

## YARN Daemons

- ResourceManager (RM)
  - Runs on master node
  - Global resource scheduler
  - Arbitrates system resources between competing applications
  - Has a pluggable scheduler to support different algorithms (such as Capacity or Fair Scheduler)
- NodeManager (NM)
  - Runs on worker nodes
  - Communicates with RM
  - Manages node resources
  - Launches containers

# Architecture Yarn Cluster (2)

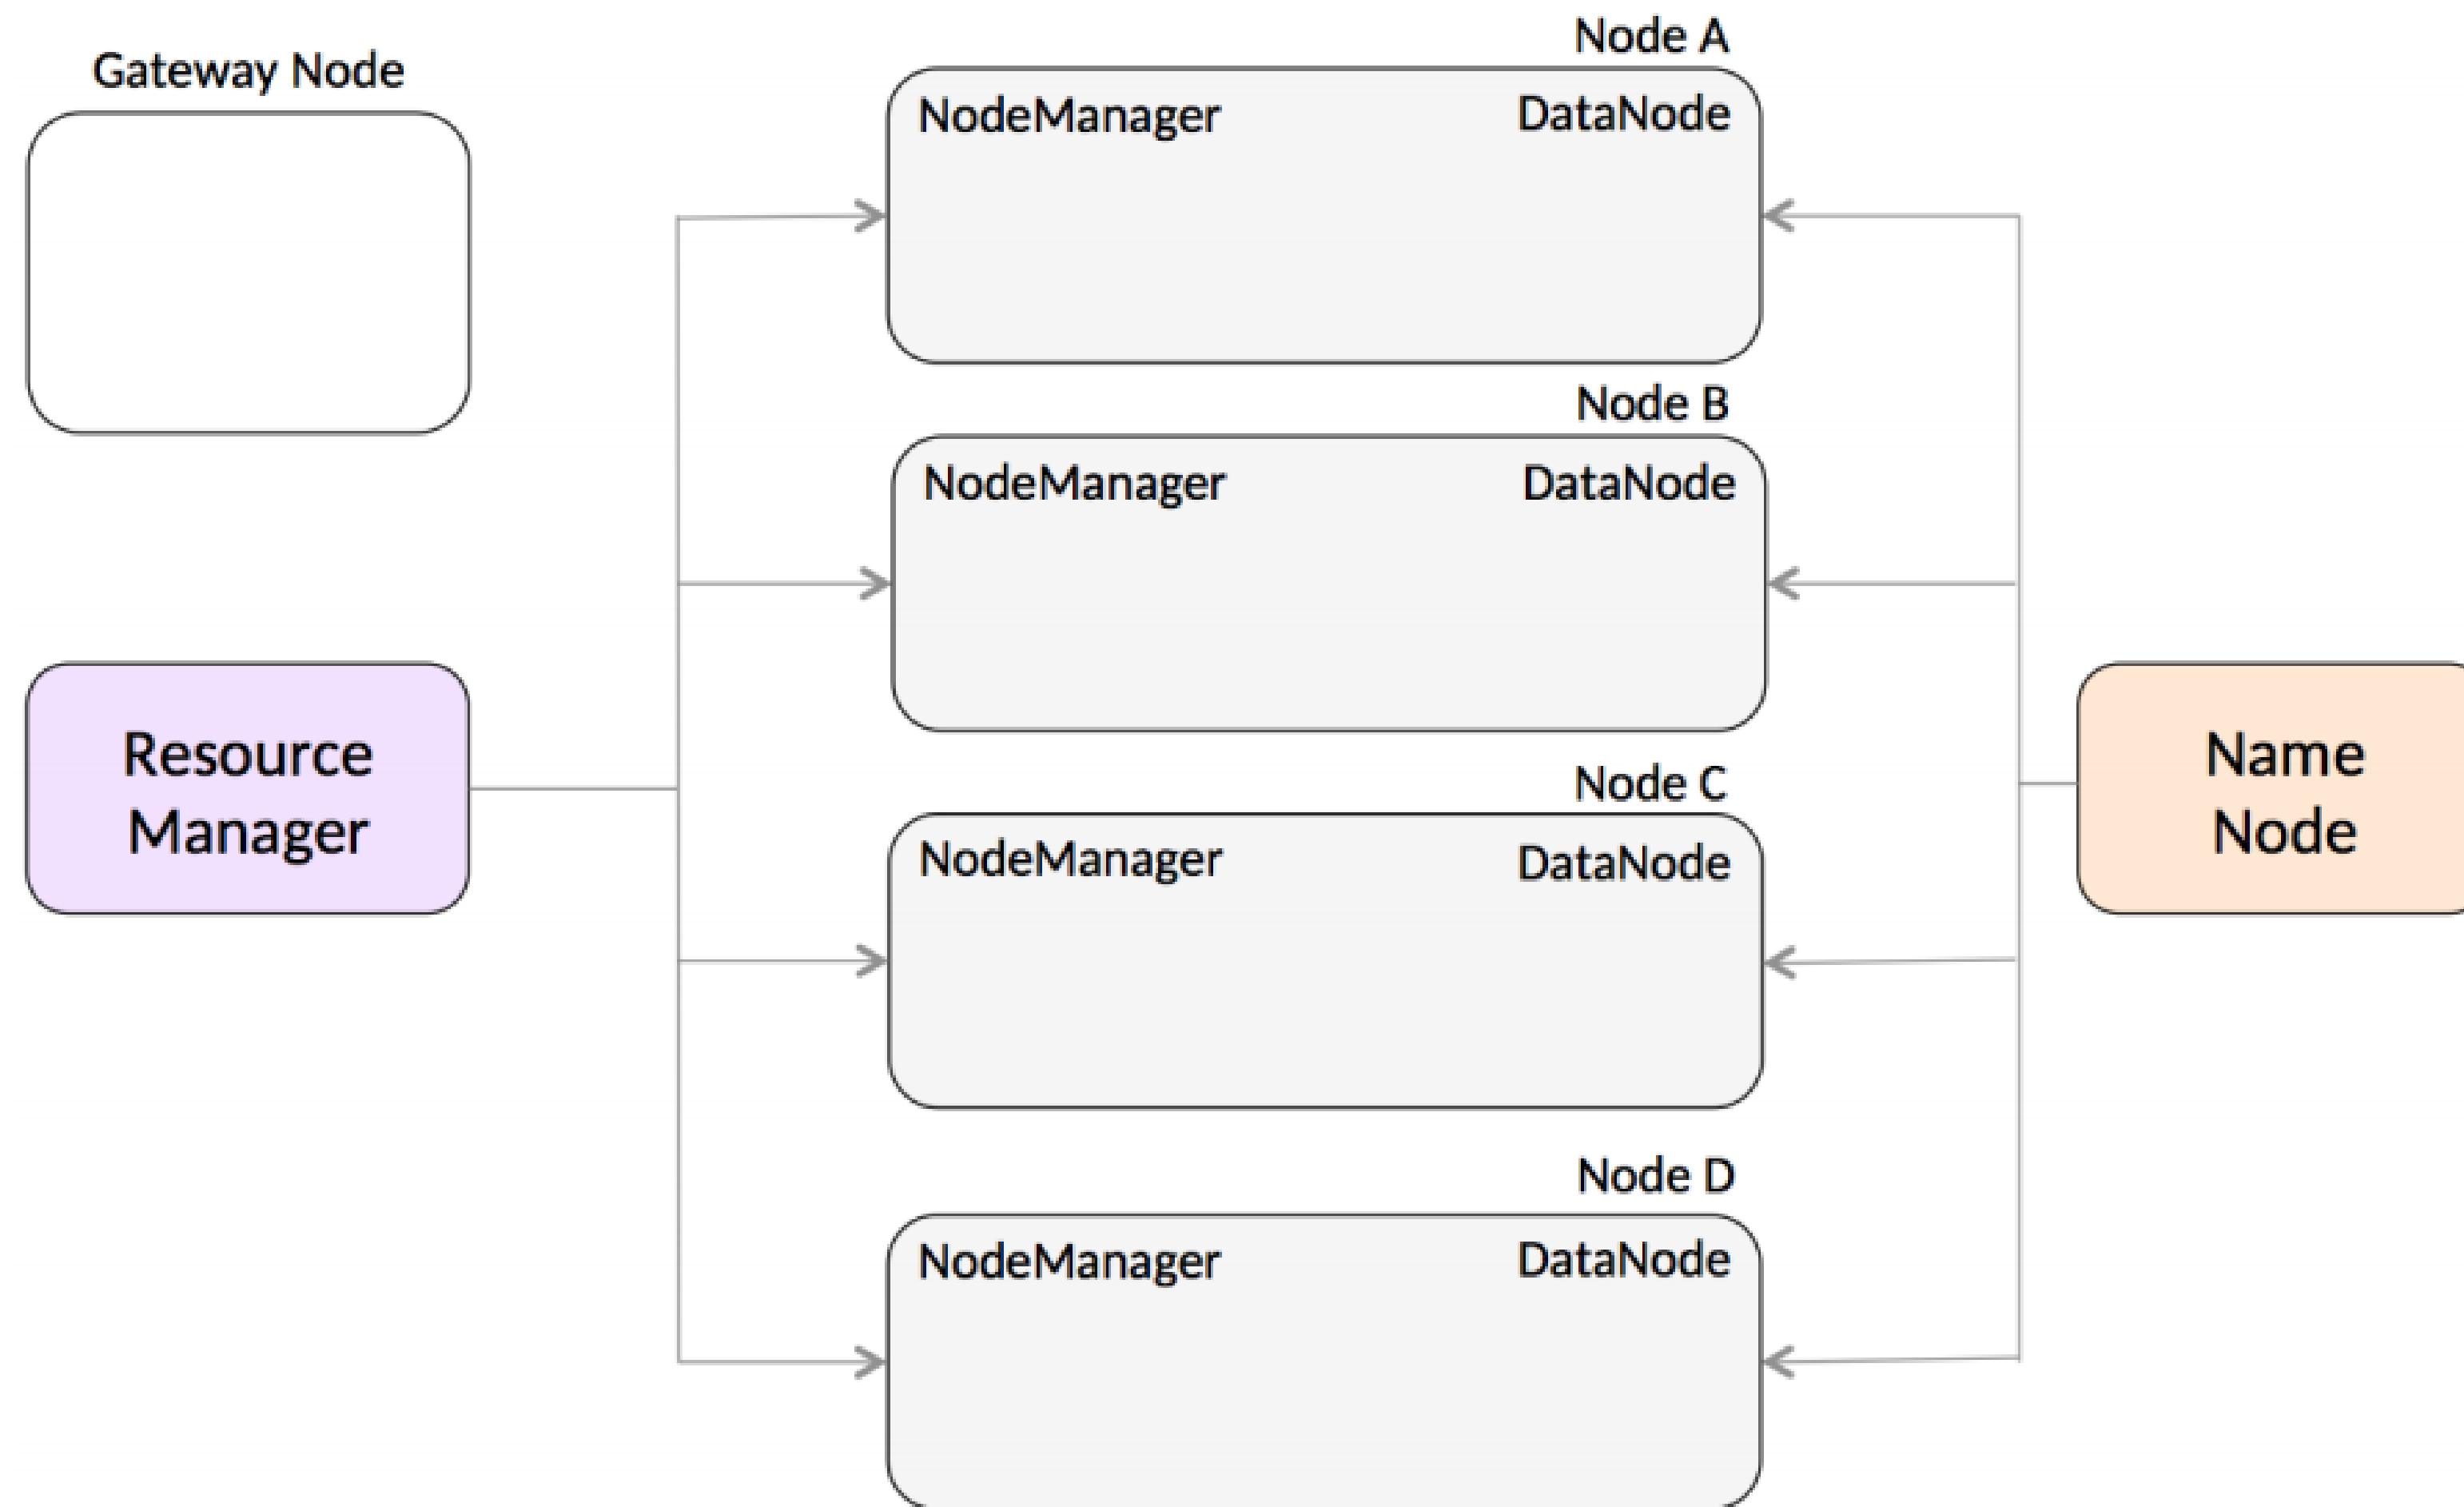
- Containers

- Containers allocate a certain amount of resources (memory, CPU cores) on a worker node
- Applications run in one or more containers
- Applications request containers from RM

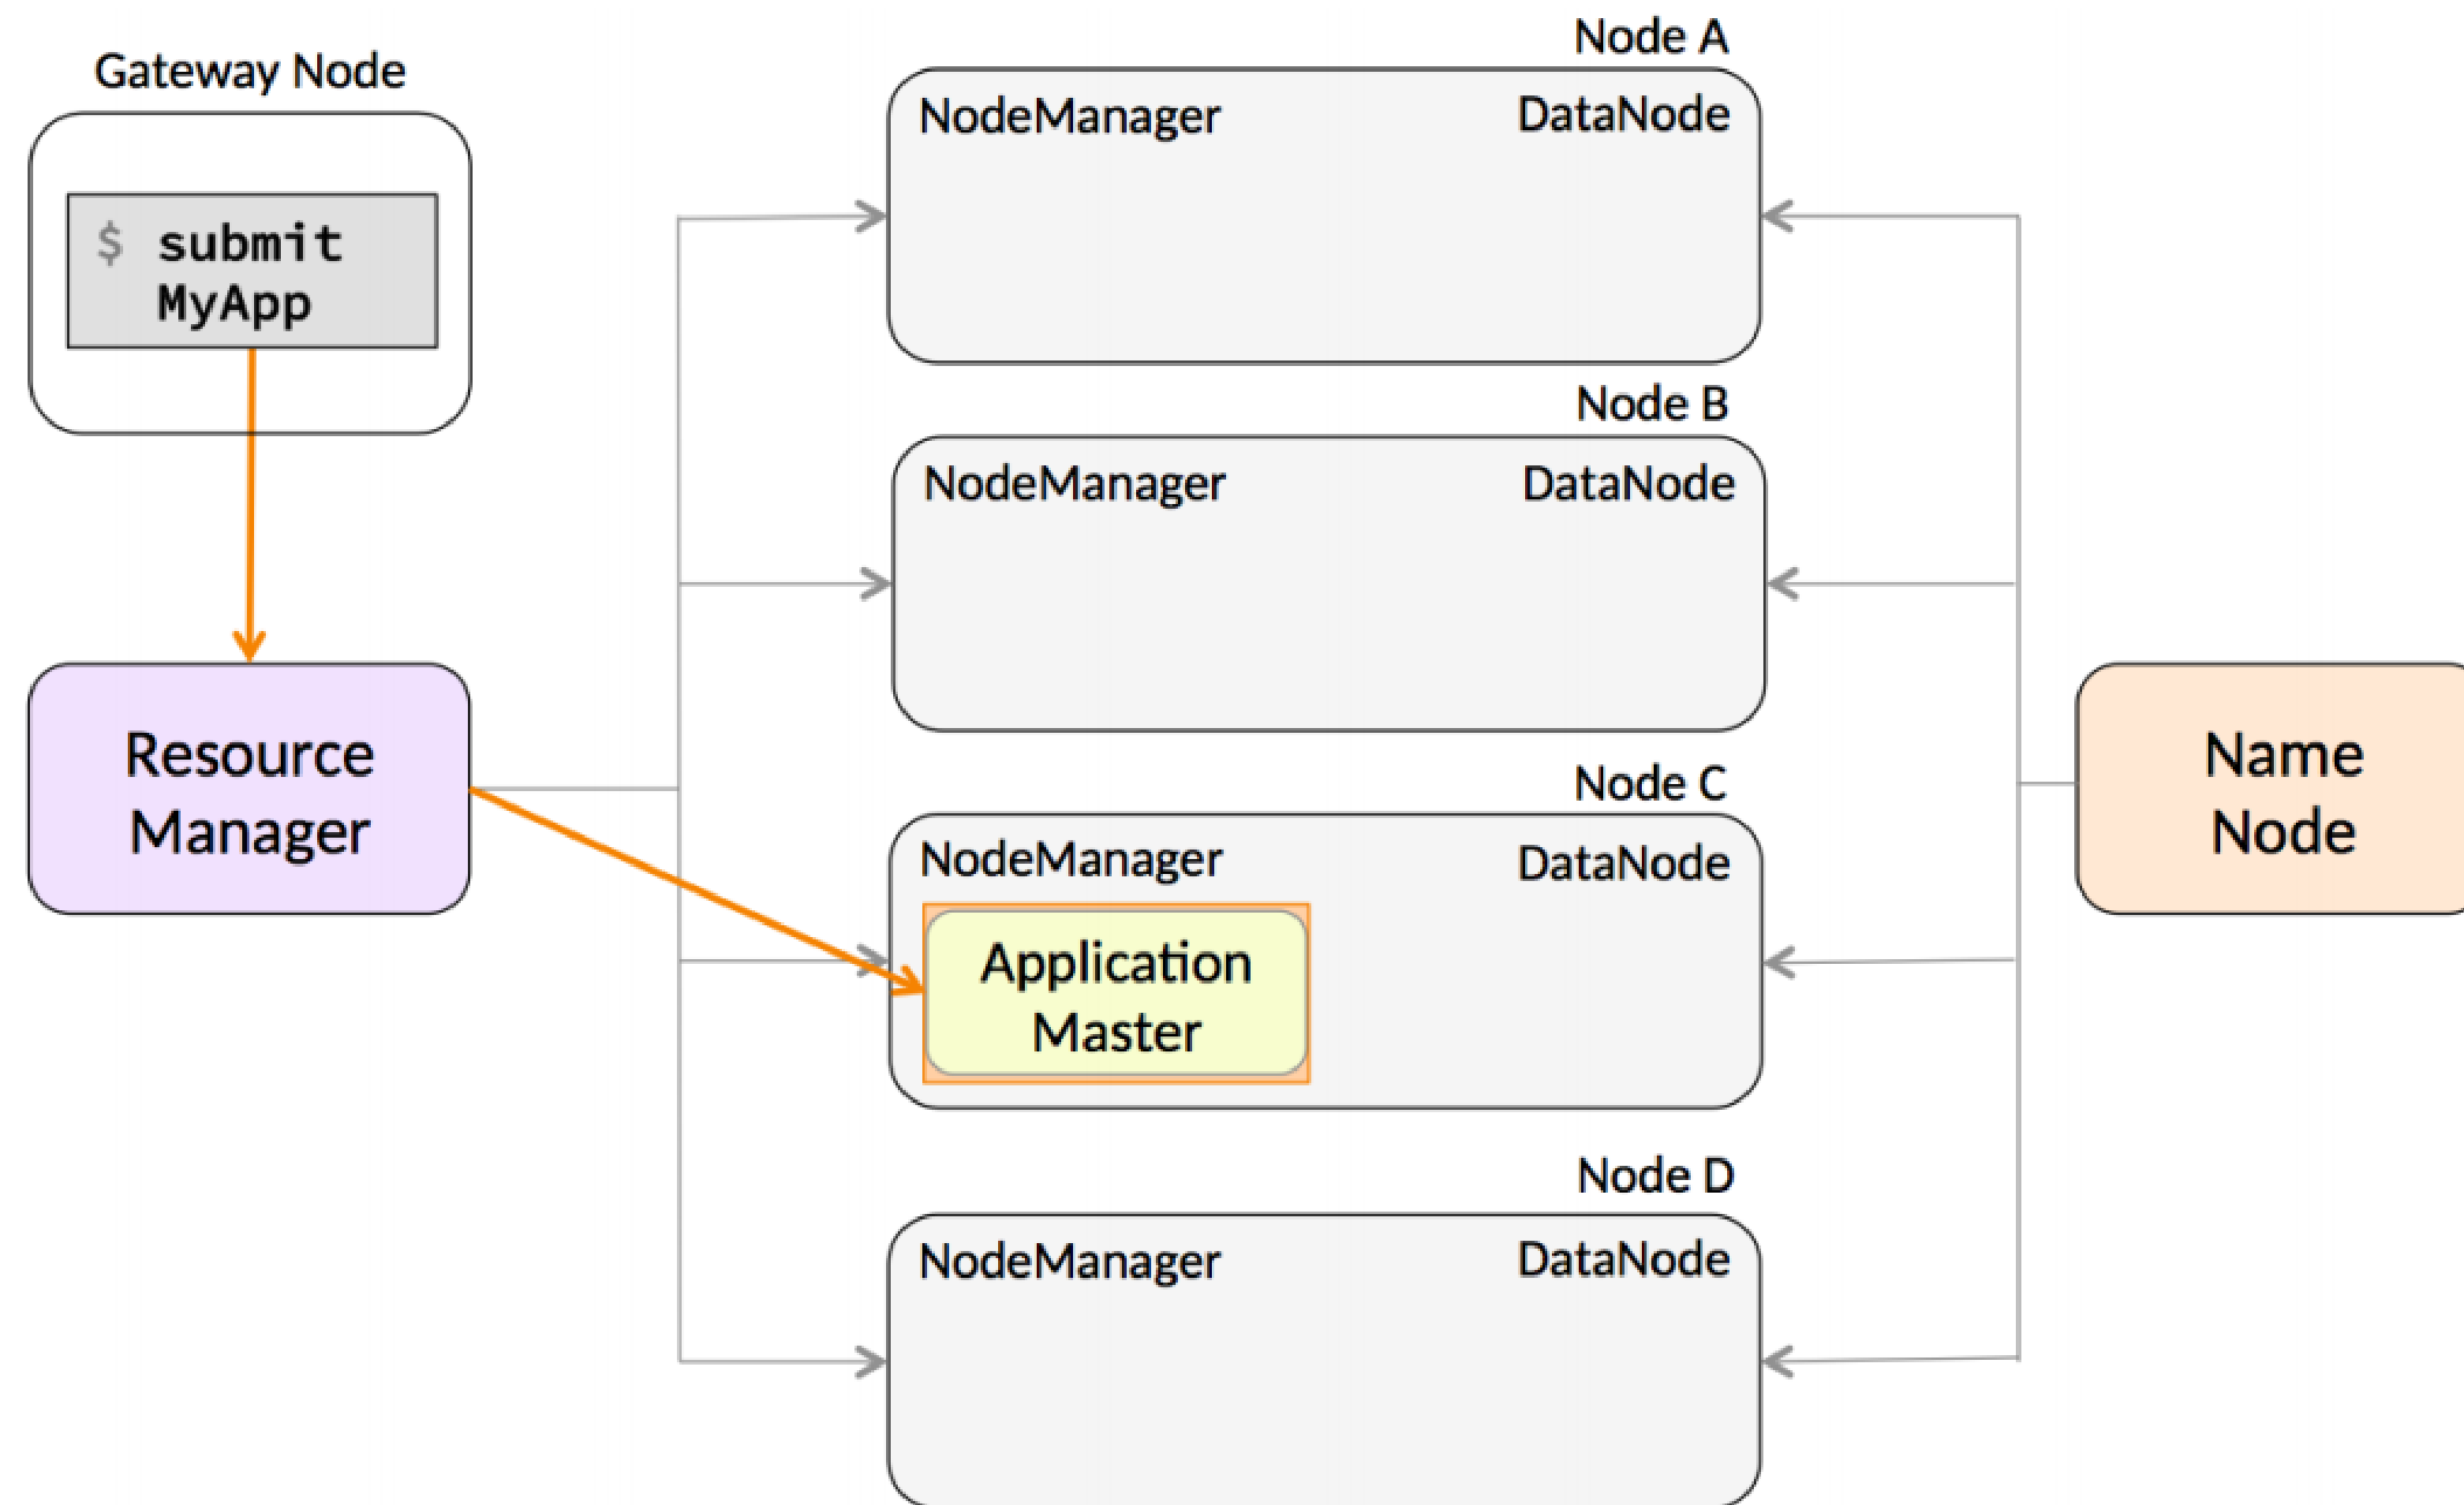
- ApplicationMaster (AM)

- One per application
- Framework/application specific
- Runs in a container
- Requests more containers to run application tasks

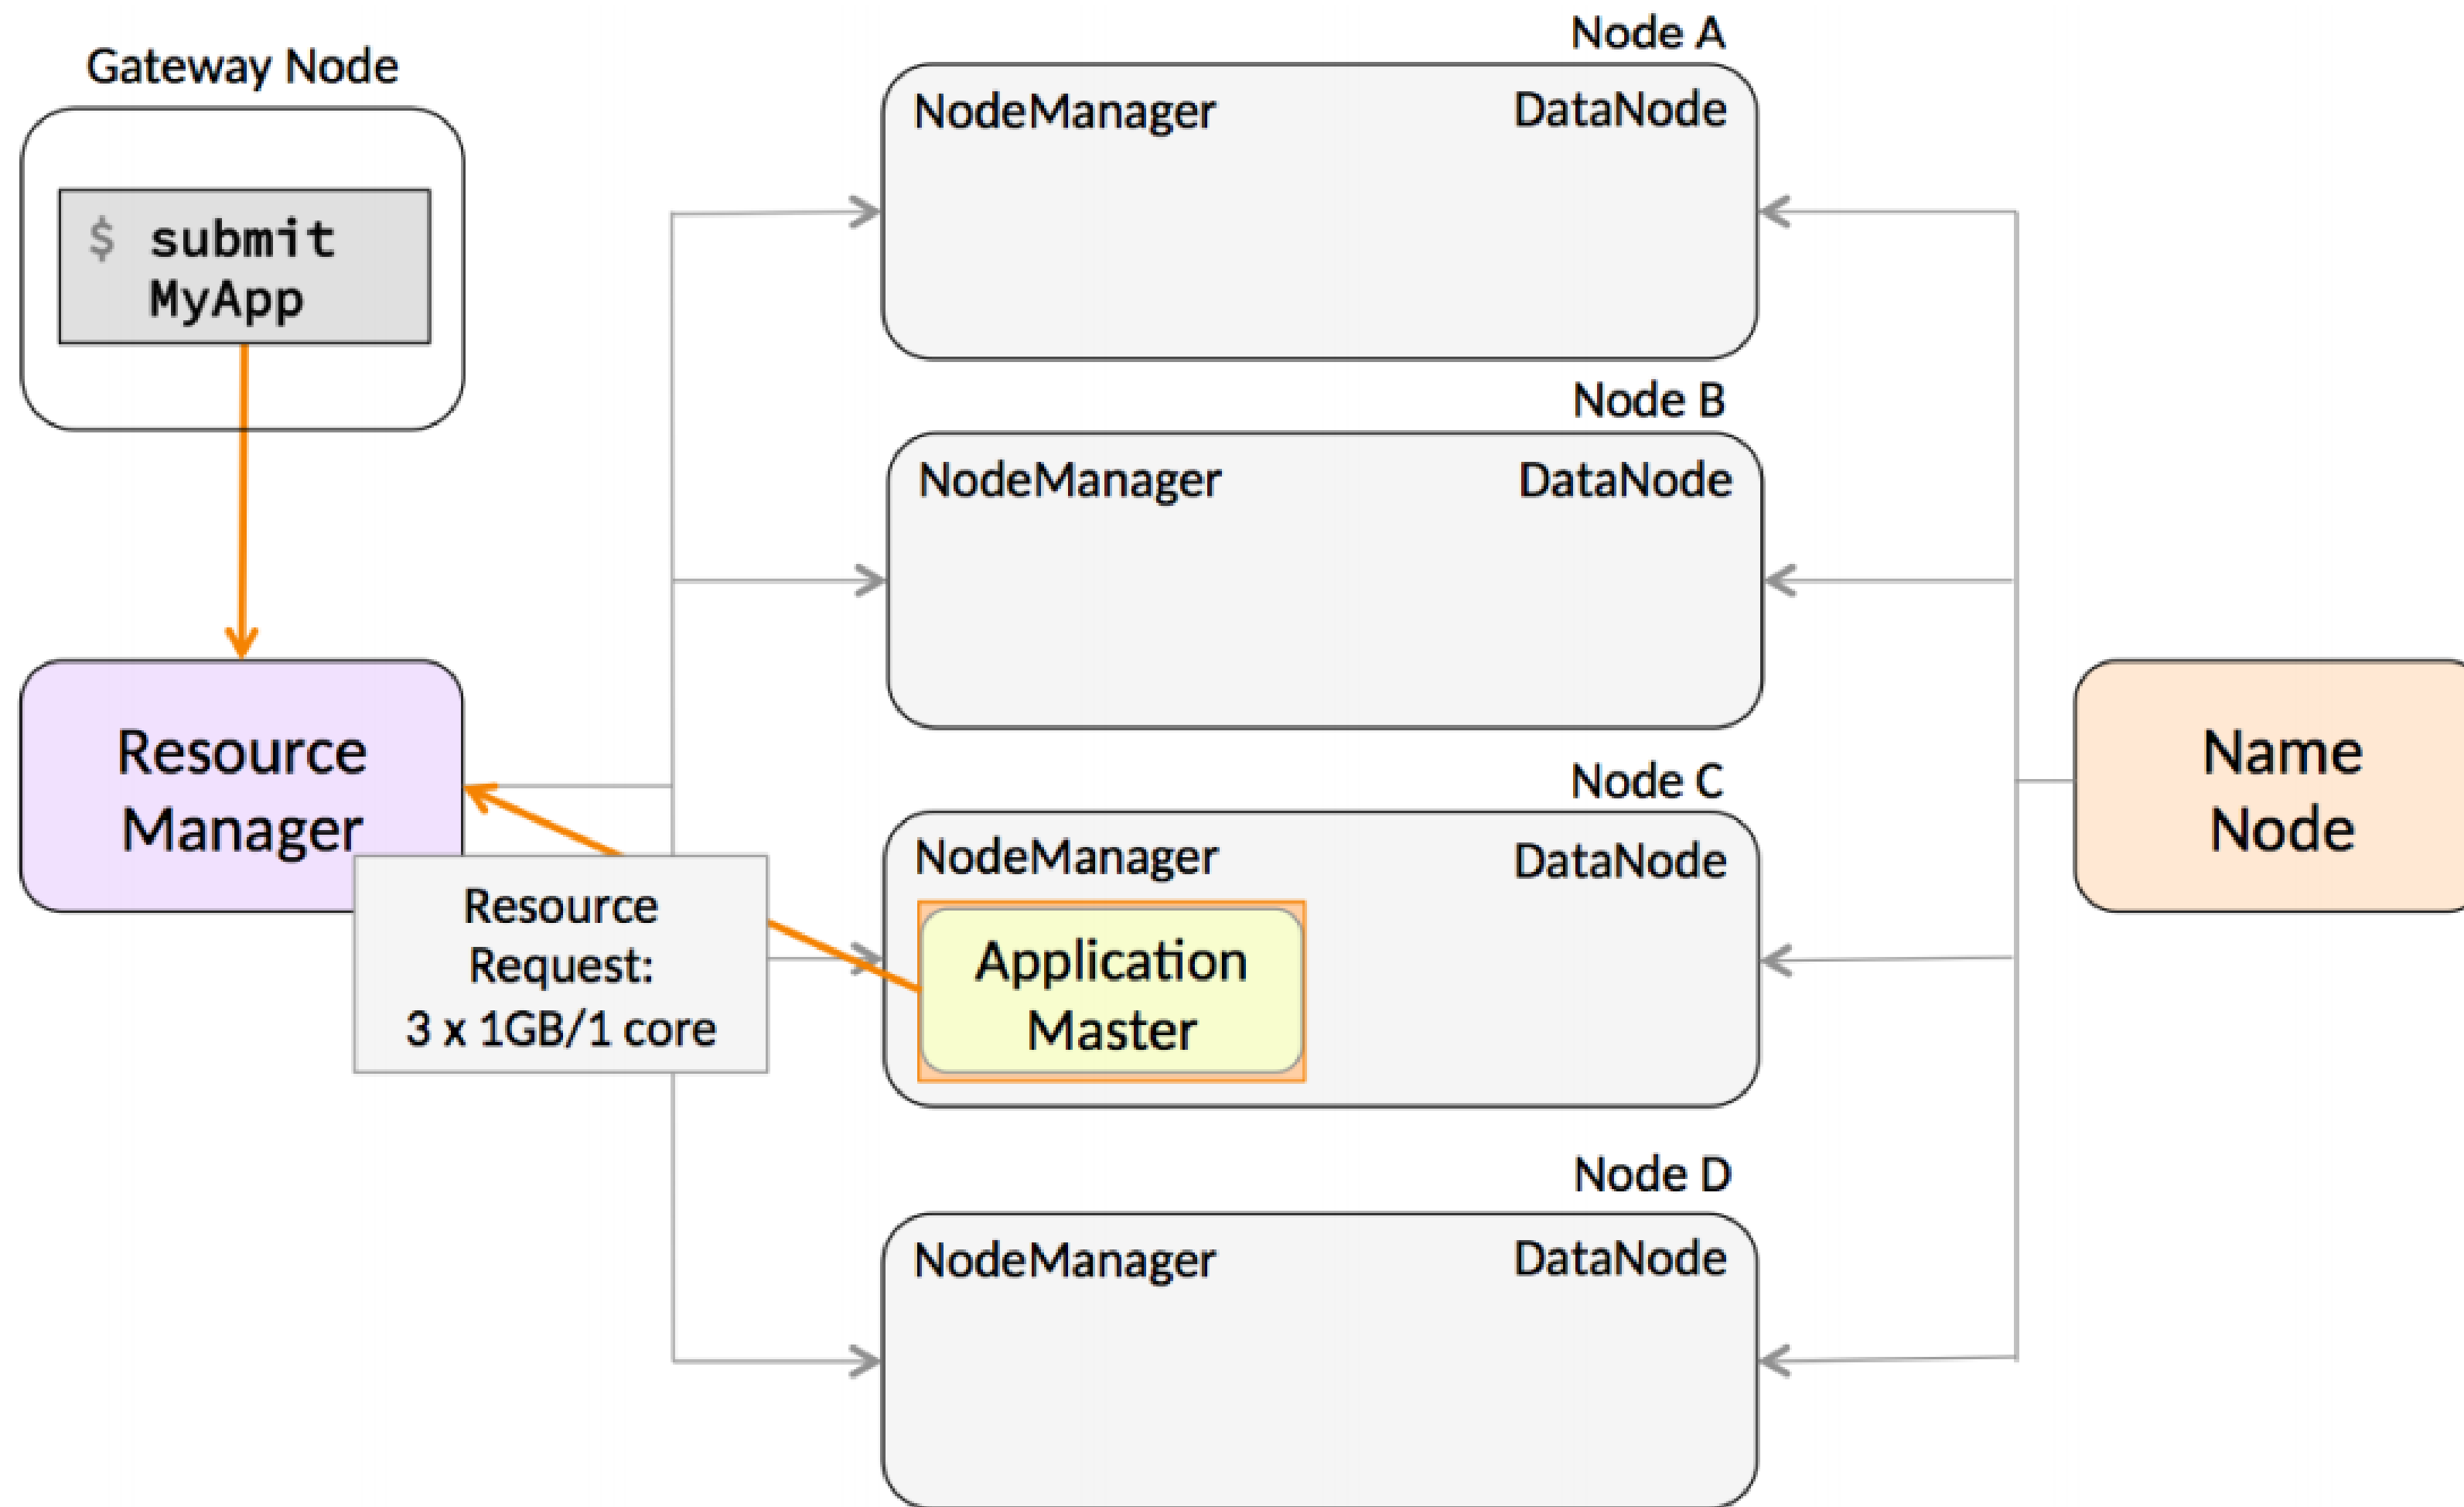
## Architecture Yarn Cluster (2)



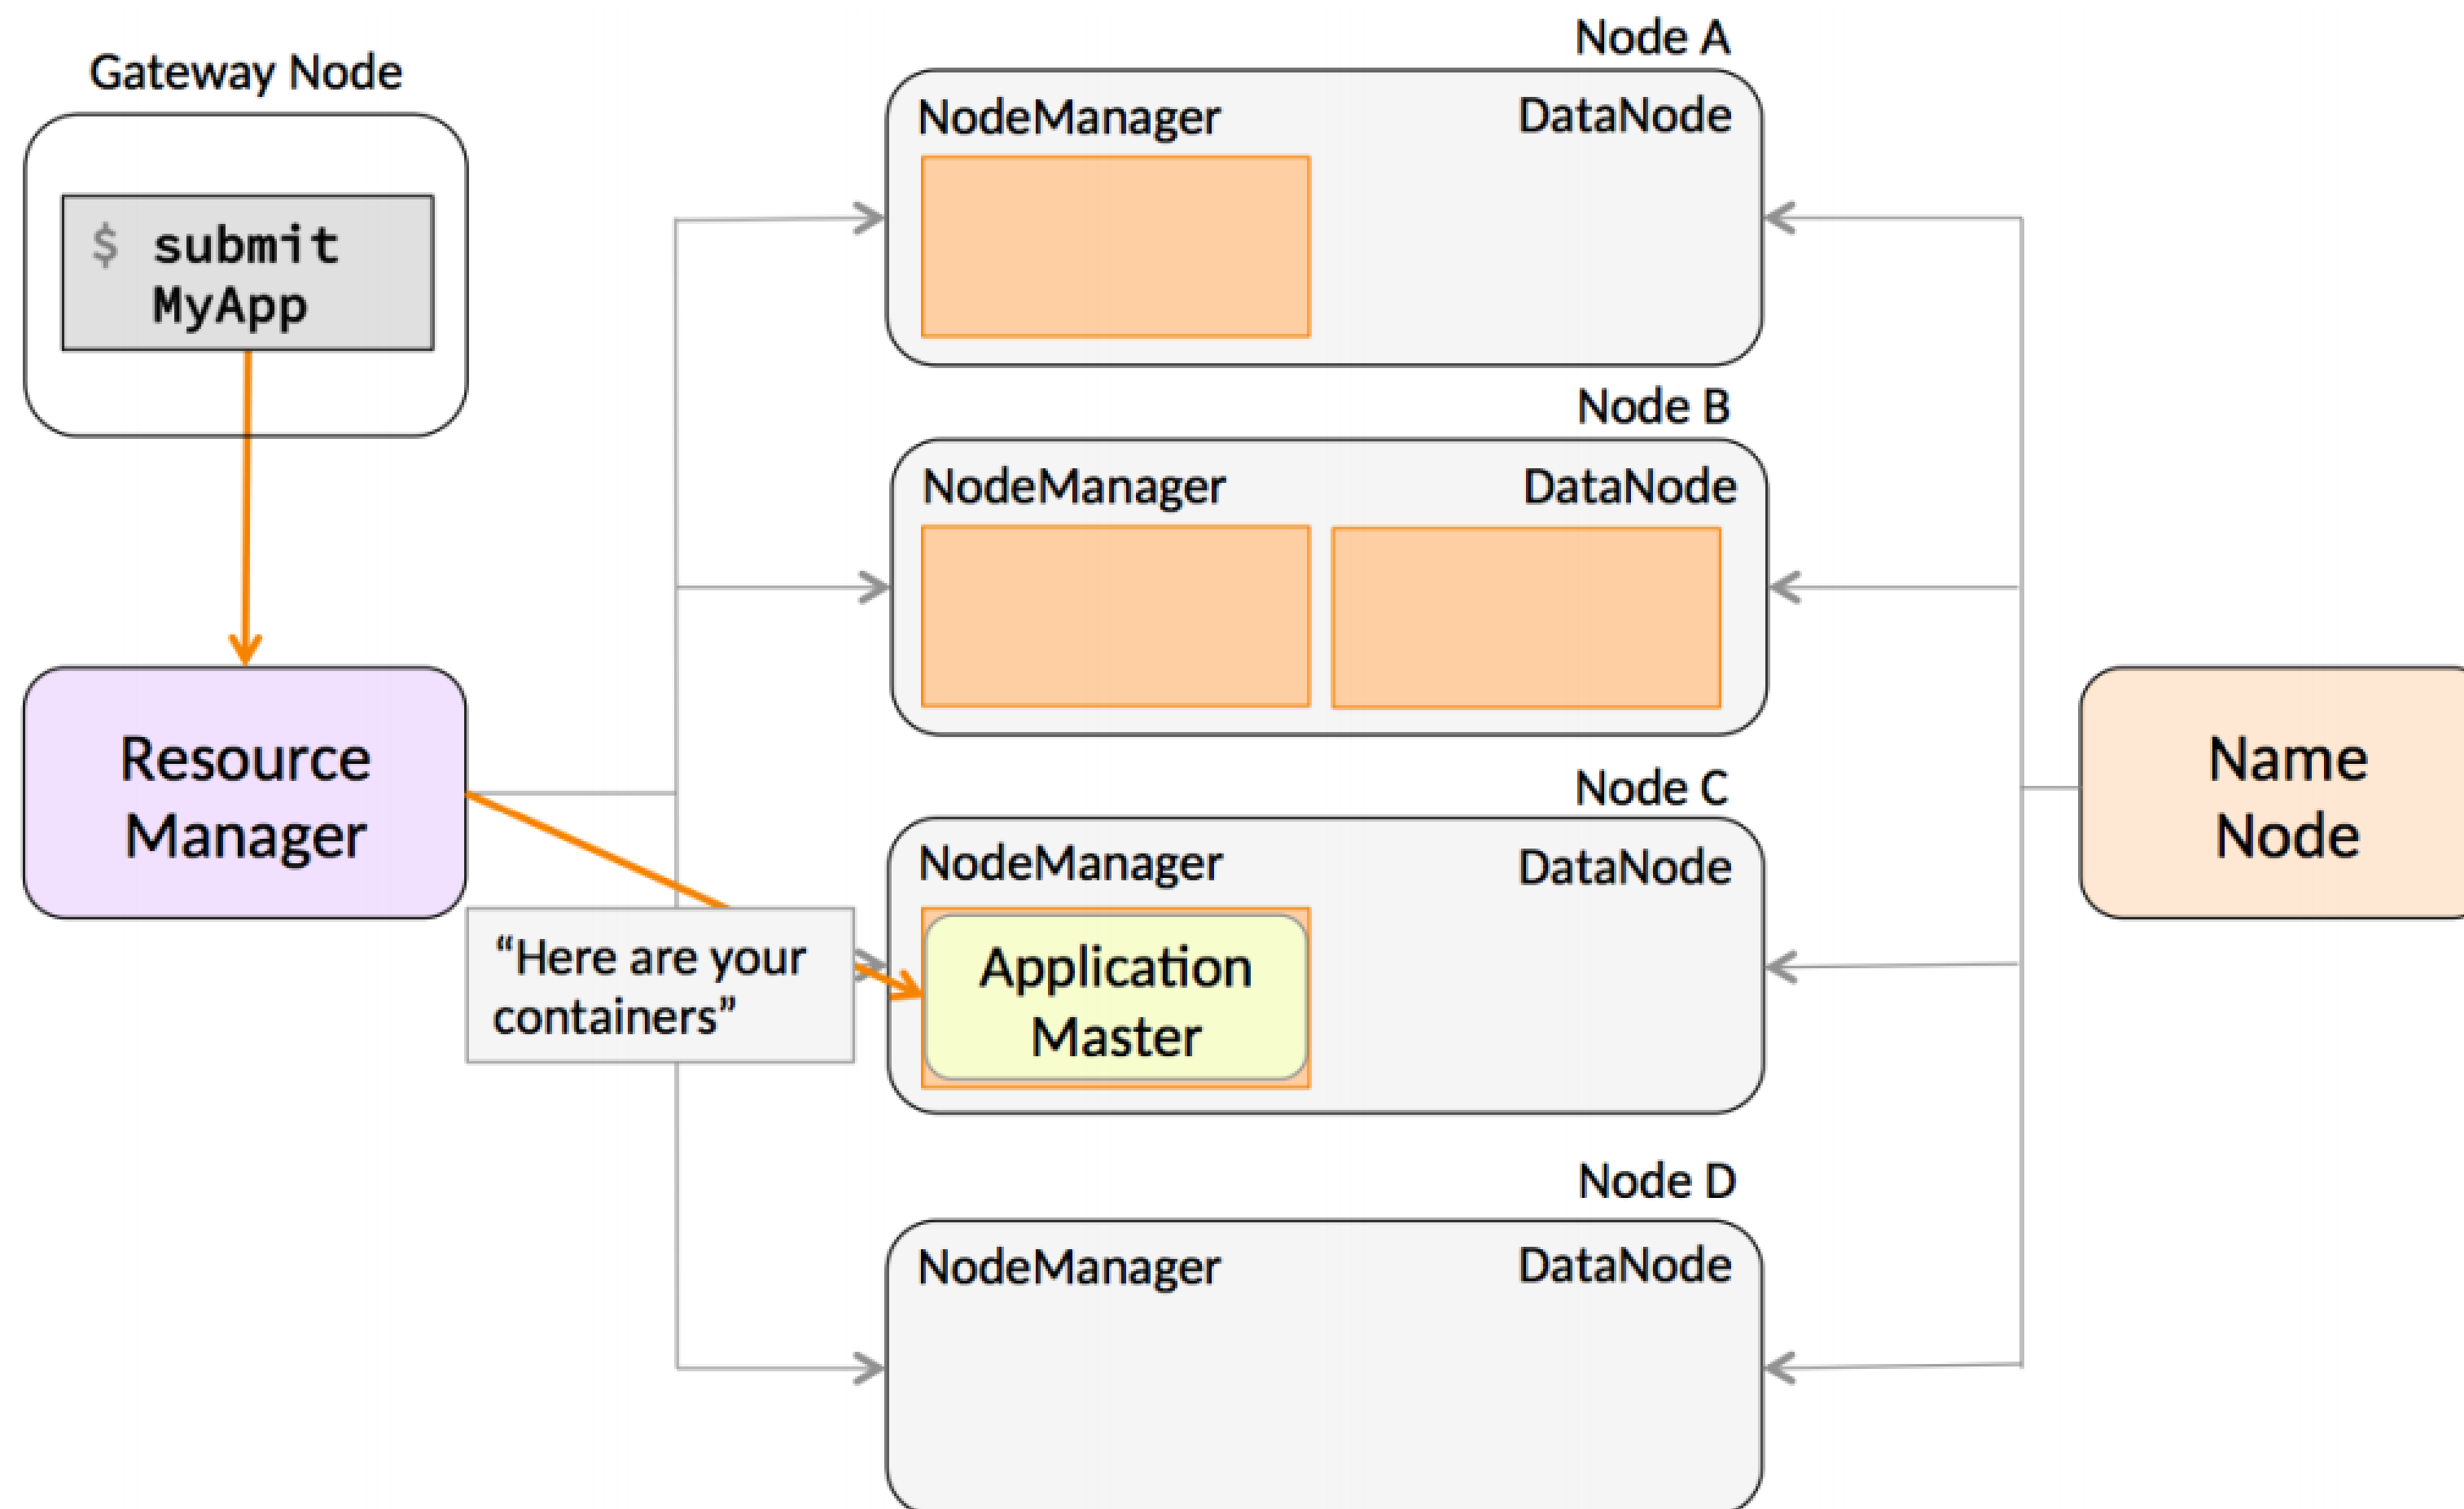
# Architecture Yarn Cluster (3)



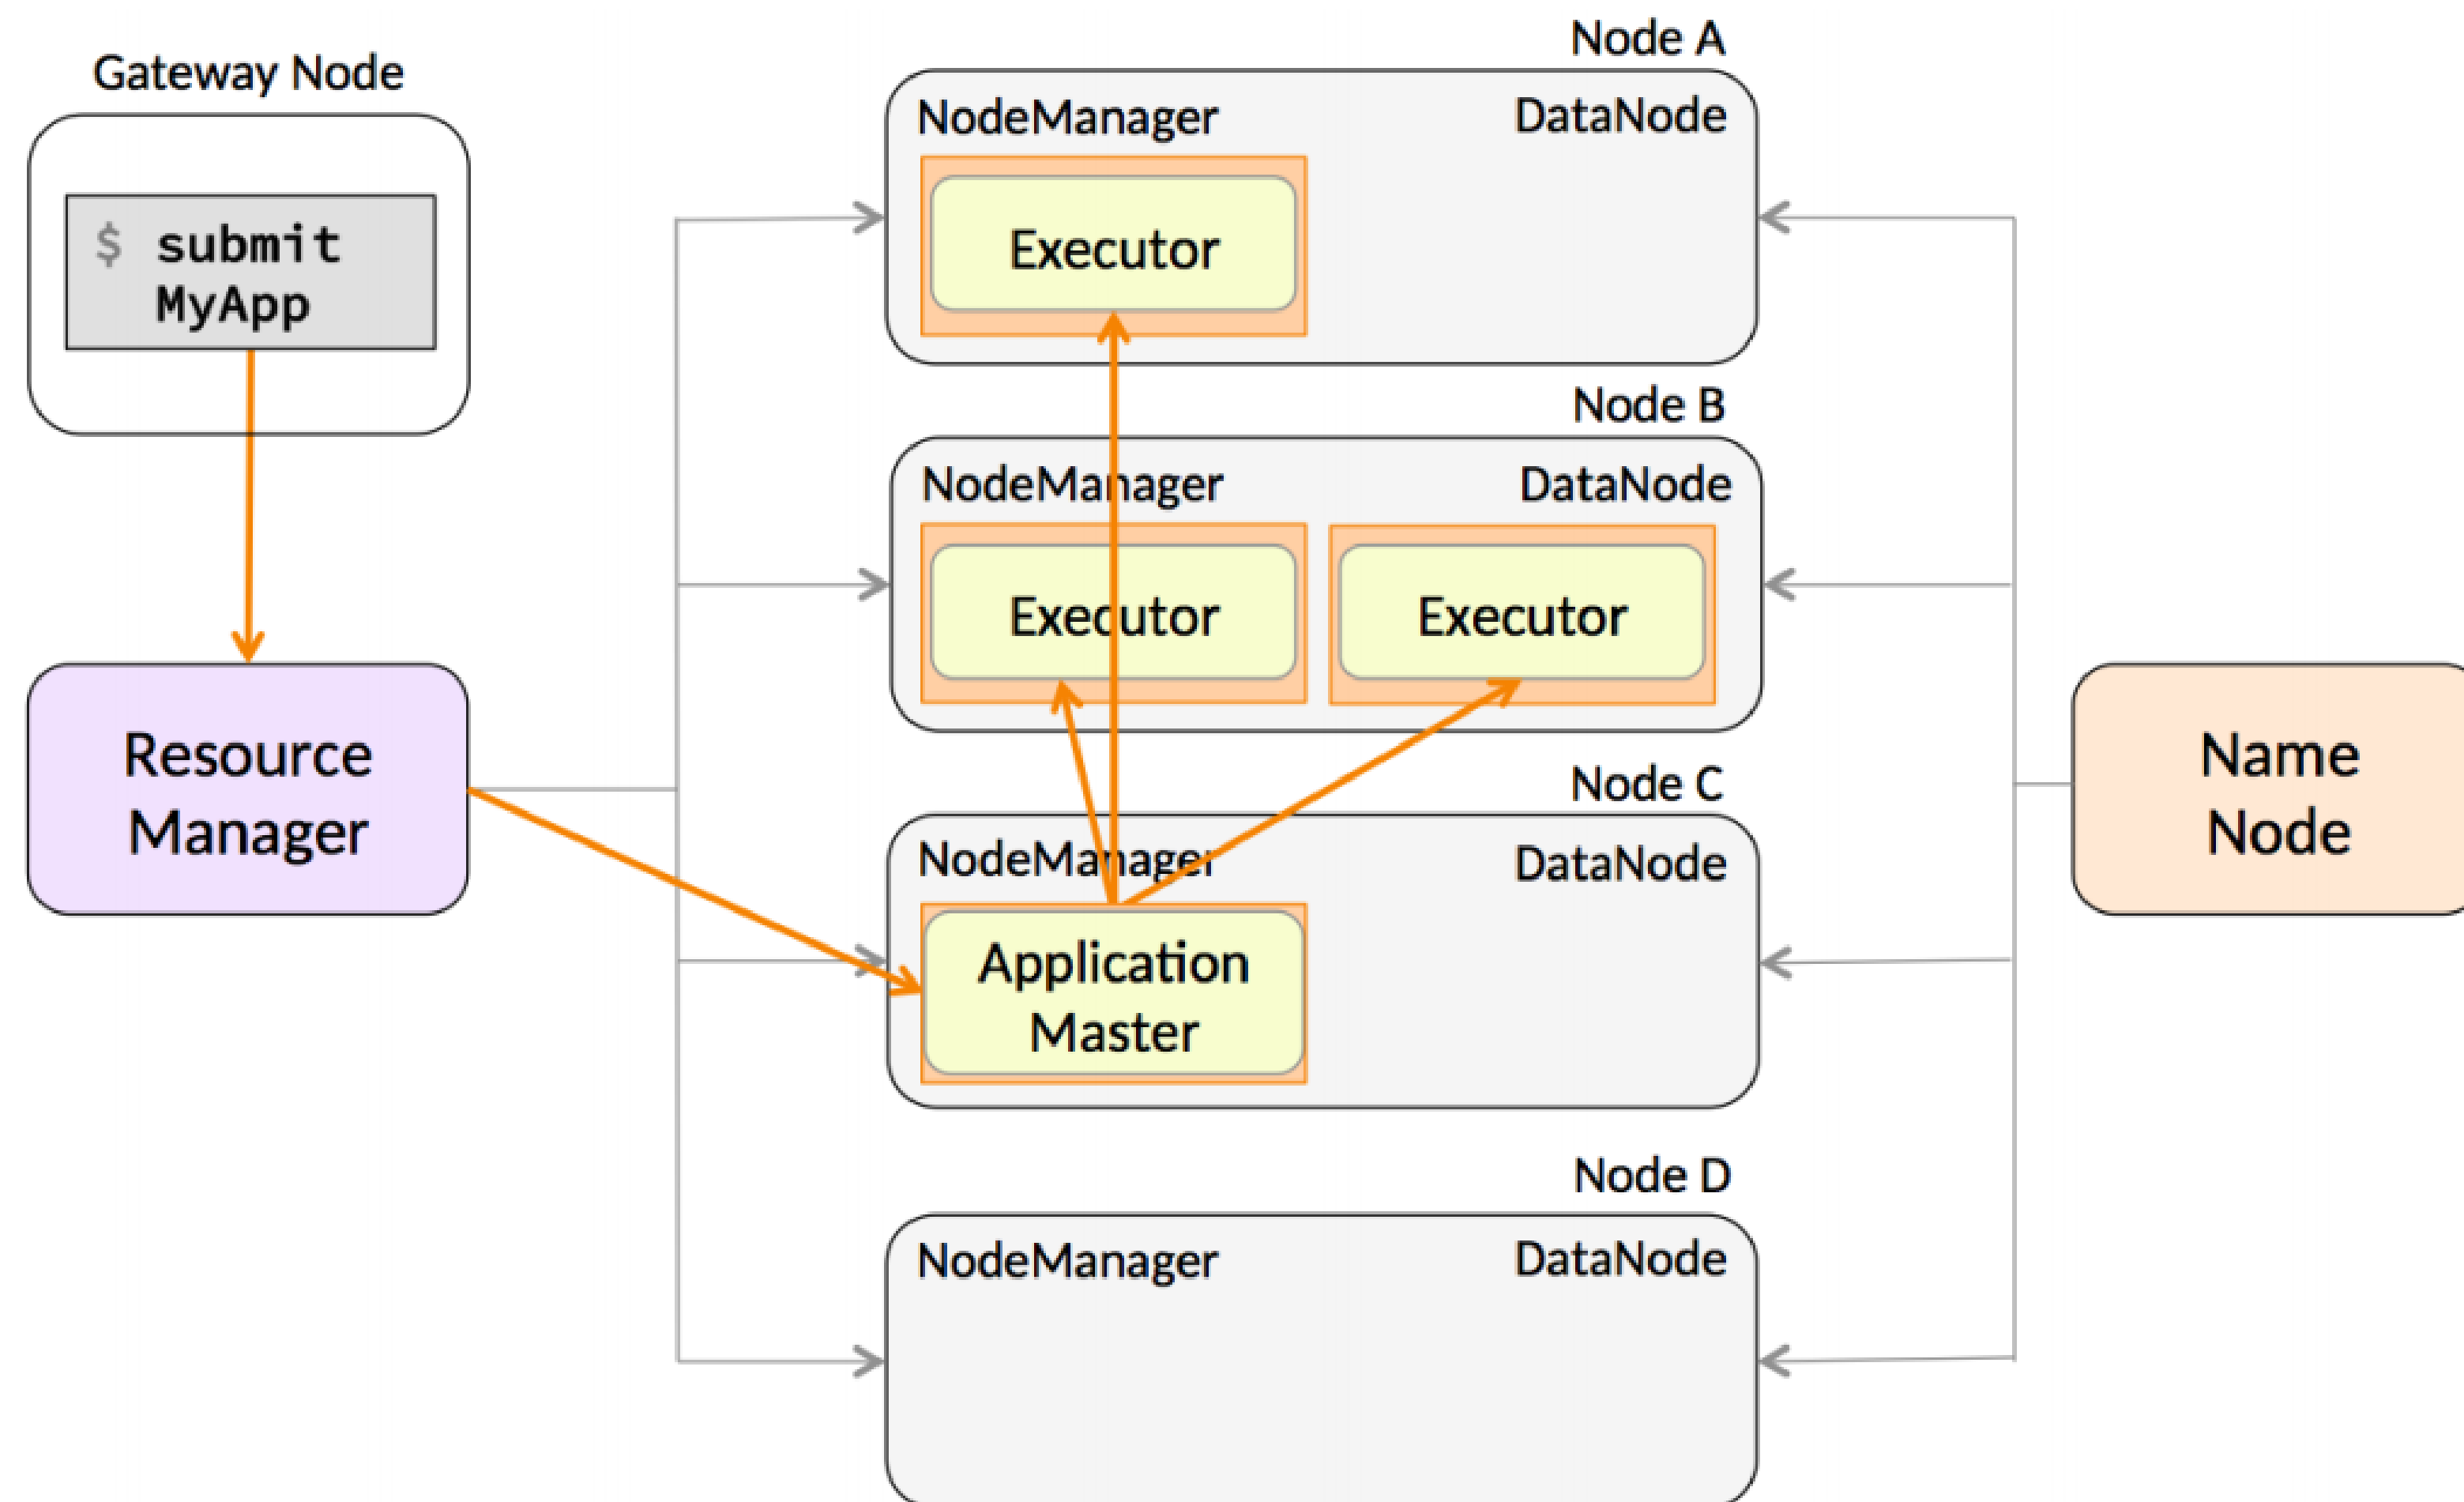
# Architecture Yarn Cluster (4)



# Architecture Yarn Cluster (5)

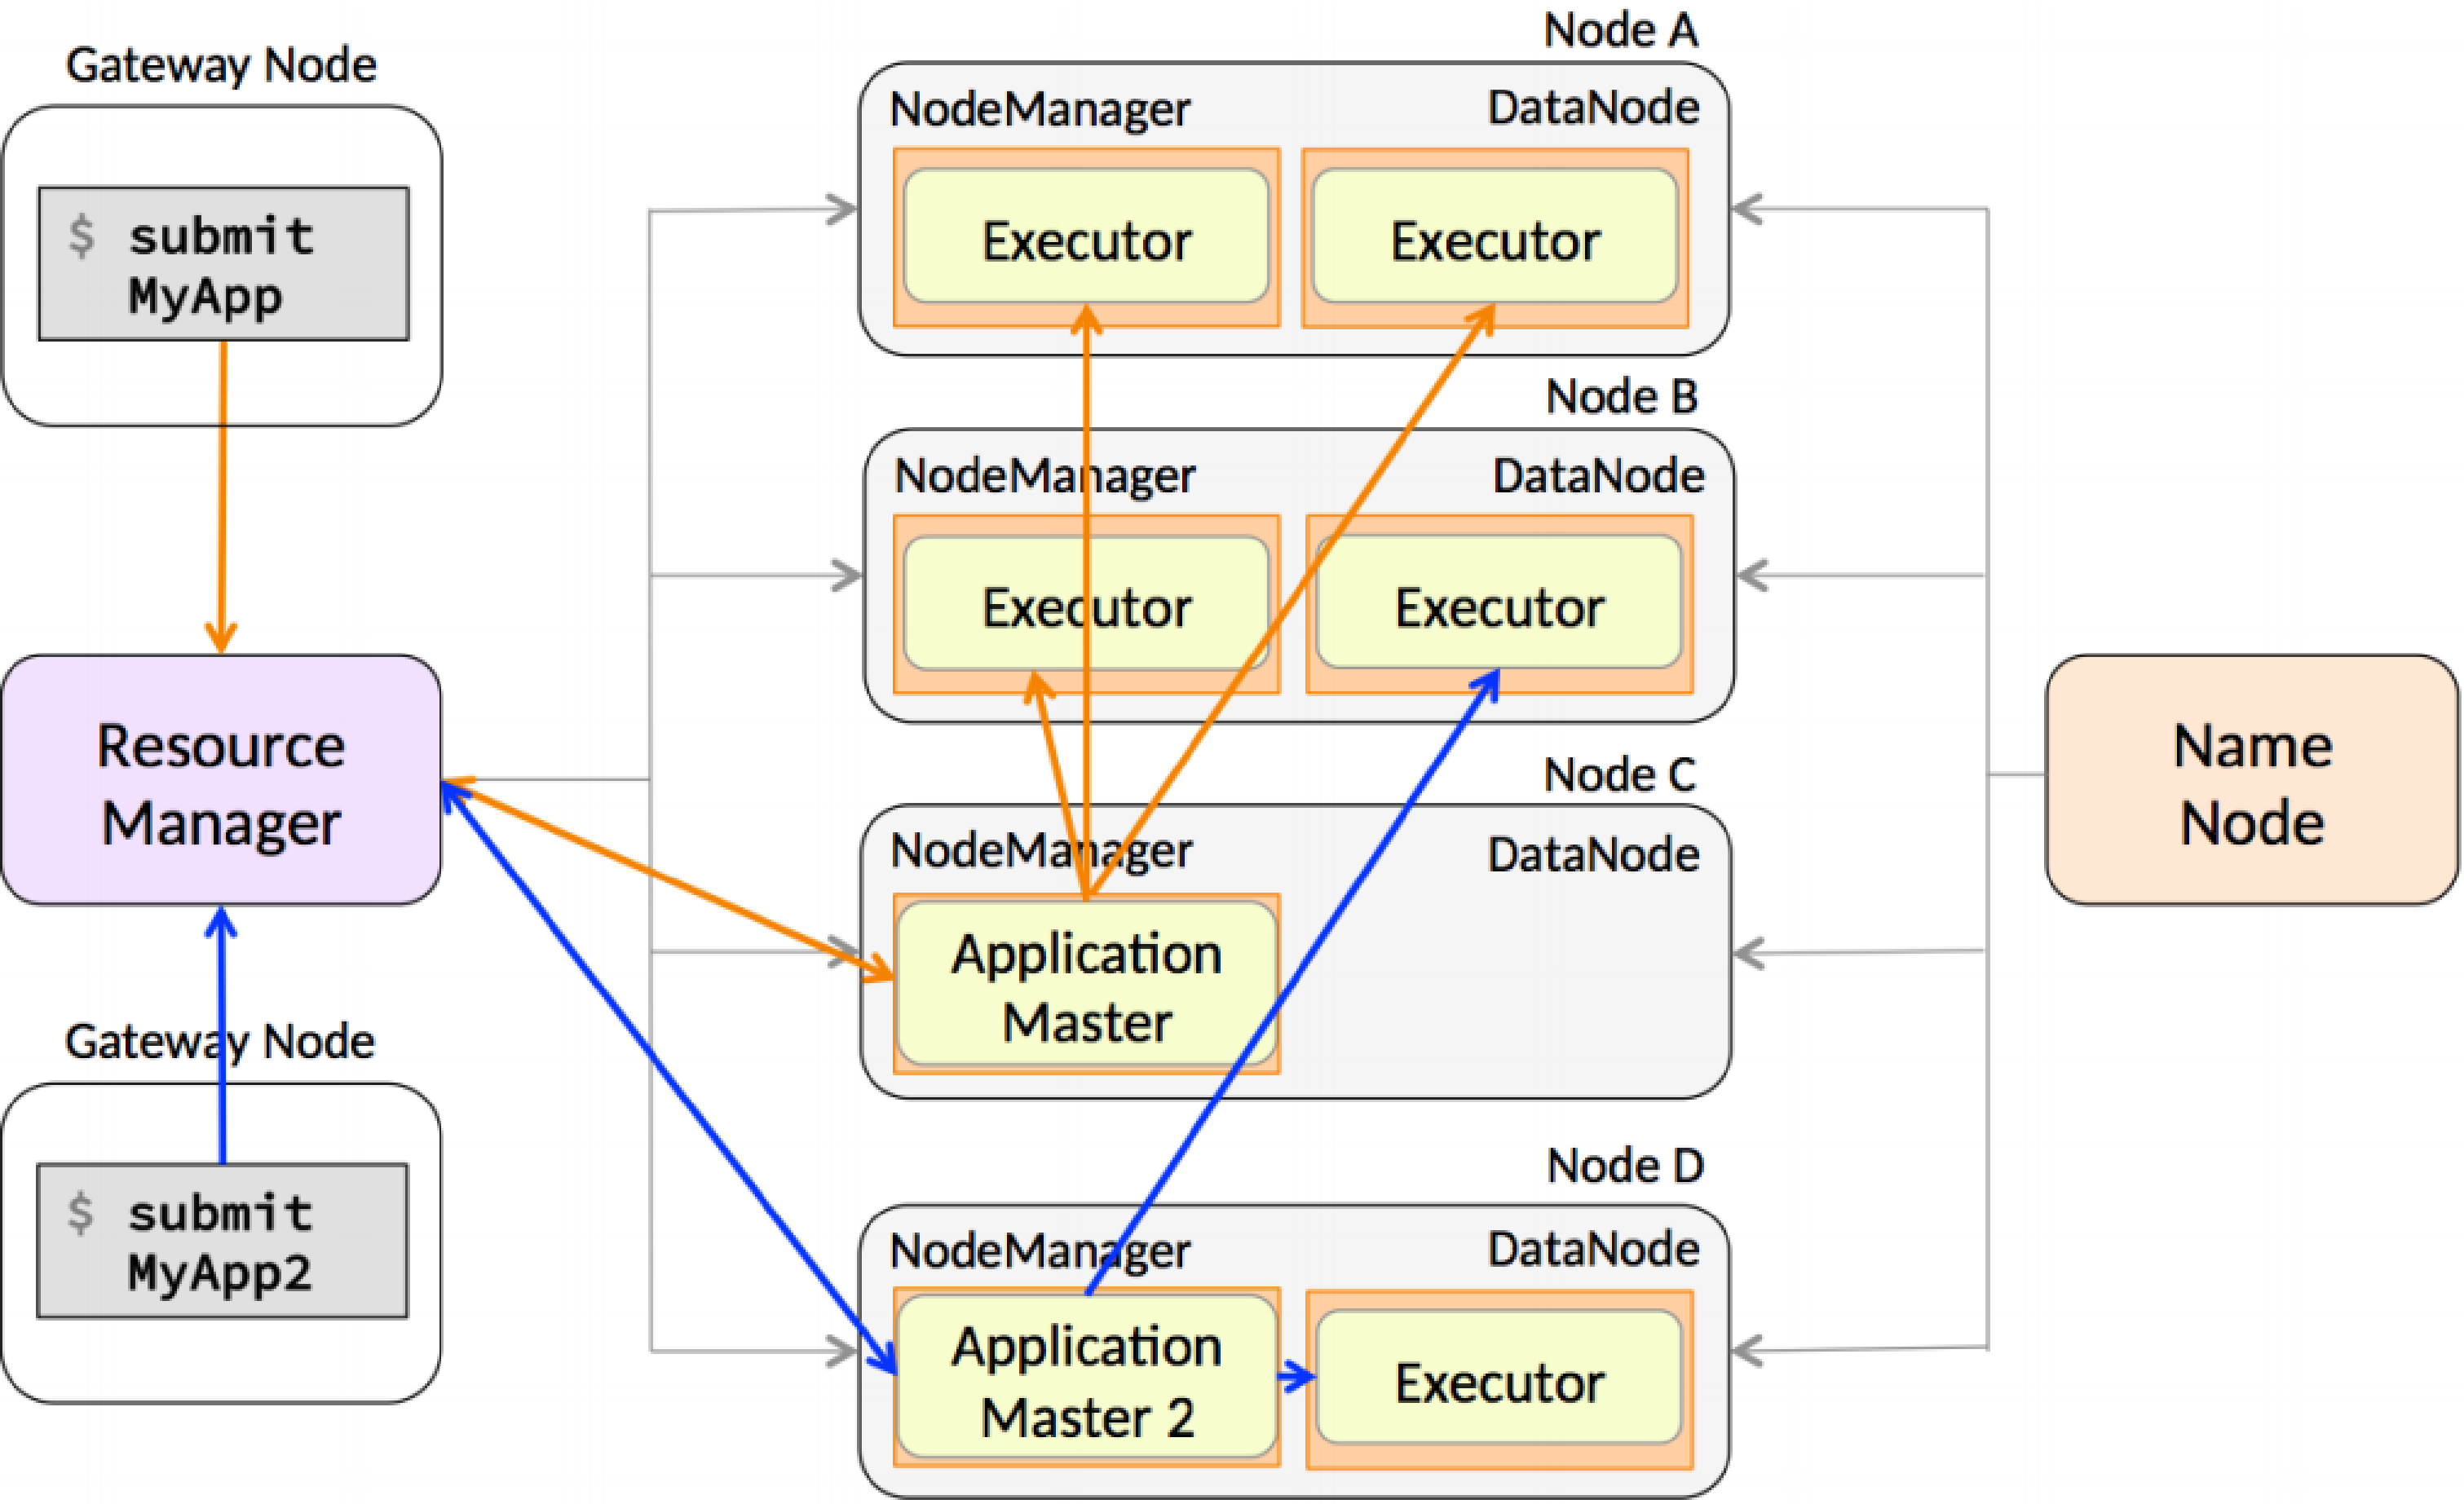


# Architecture Yarn Cluster (6)





# Architecture Yarn Cluster (7)



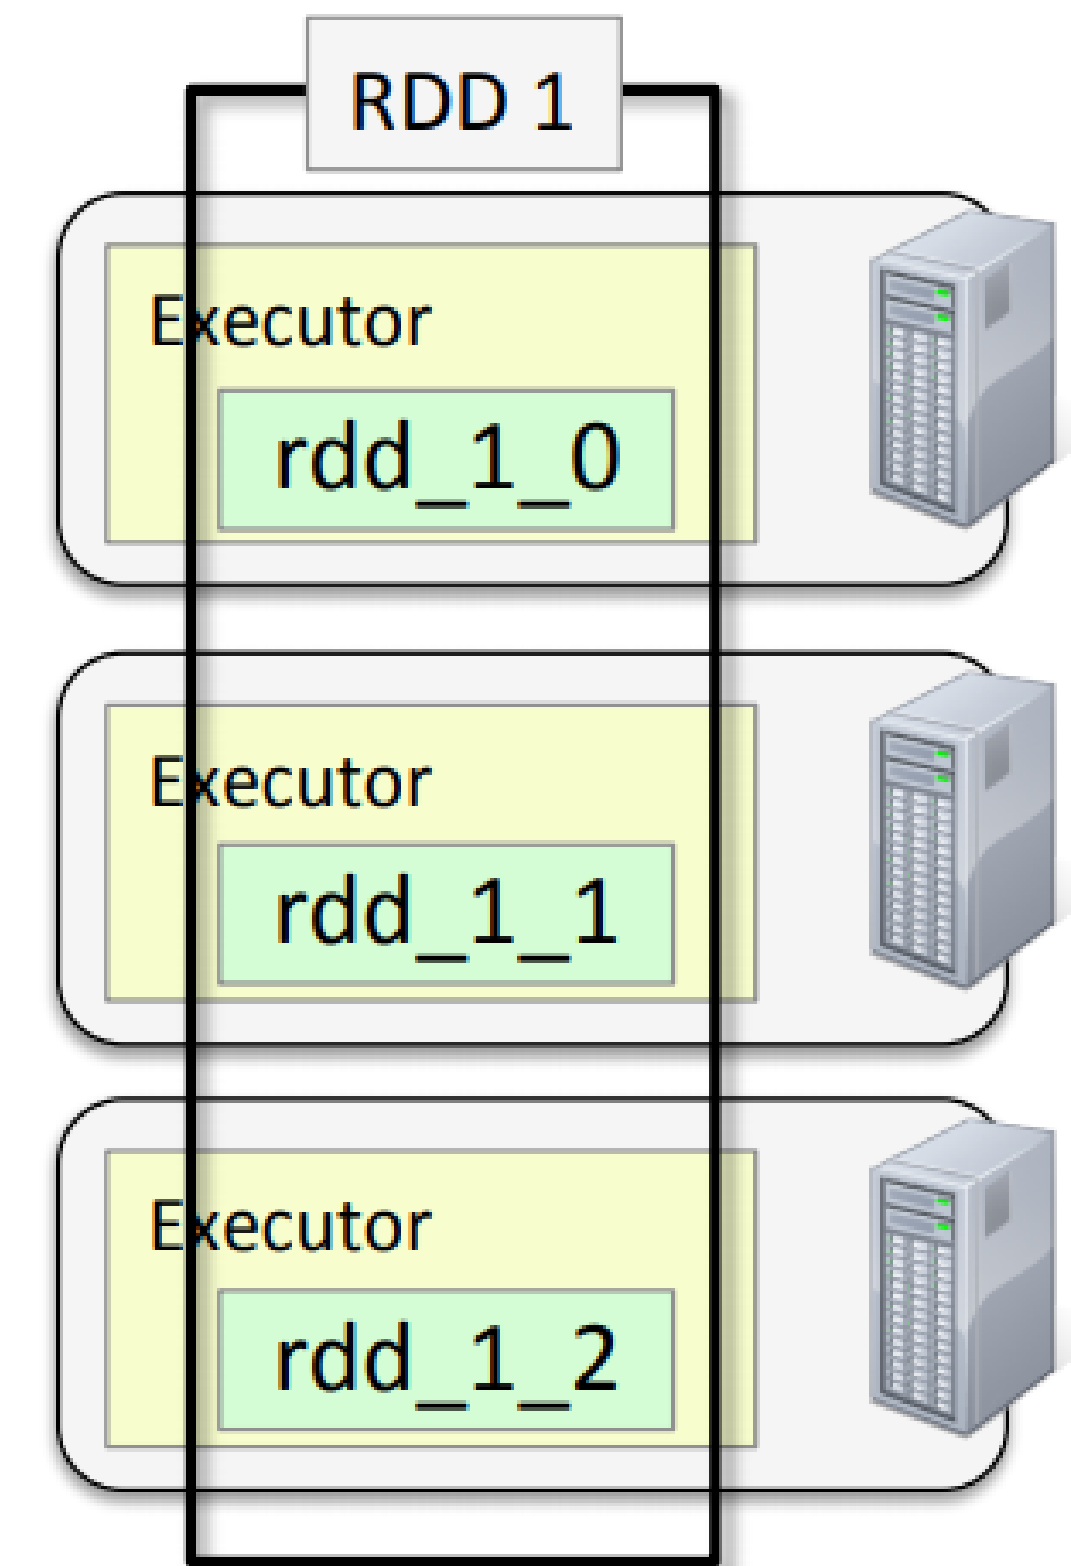
# Data model: Resilient Distributed Datasets (RDDs)(1)

- RDDs are part of core Spark
- Resilient Distributed Dataset (RDD)
  - Resilient: If data in memory is lost, it can be recreated
  - Distributed: Processed across the cluster
  - Dataset: Initial data can come from a source such as a file, or it can be created programmatically
- RDDs are unstructured
  - No schema defining columns and rows
  - Not table-like; cannot be queried using SQL-like transformations such as where and select
  - RDD transformations use lambda functions



# Data model: Resilient Distributed Datasets (RDDs)(1)

- Resilient Distributed Datasets
  - Data is partitioned across worker nodes
- Partitioning is done automatically by Spark
  - Optionally you can control how many partitions are created



# Data model: DataFrame

- DataFrames represent structured data in a tabular form
  - DataFrames model data similar to tables in an RDBMS
  - DataFrames consist of a collection of loosely typed Row objects
  - Rows are organized into columns described by a schema
- DataFrames contain an ordered collection of Row objects
  - Rows contain an ordered collection of values
  - Row values can be basic types (such as integers, strings, and floats) or collections of those types (such as arrays and lists)
  - A schema maps column names and types to the values in a row



# Data model: DataFrame

- DataSet is distributed collection of strongly-typed objects
  - Primitive types such as Int or String
  - Complex types such as arrays and lists containing supported types
- Mapped to a relational schema
  - The schema is defined by an encoder
  - The schema maps object properties to typed columns
- In Scala, DataFrame is an alias for a Dataset containing Row objects
- DataFrames and Datasets represent different types of data
  - DataFrames (Datasets of Row objects) represent tabular data
  - — Datasets represent typed, object-oriented data



## Data model: recap

- RDD Is building block of spark. No matter which abstraction Dataframe or Dataset we use, internally final computation is done on RDDs
- offers huge performance improvement over RDDs
  - Custom Memory management
  - Optimized Execution Plans
  - but Lack of Type Safety
- DataSet is an extension to Dataframe API
  - comes with OOPs style and developer friendly compile time safety



# Data model: case uses

- RDD
  - low-level transformation
  - unstructured data
- DataFrame or DataSets
  - want rich semantics or high-level abstractions
  - need high-level expressions
  - want higher degree of type-safety at compile time

## Example: Wordcount(1)

Input Data

the cat sat on the mat  
the aardvark sat on the sofa



Result

aardvark	1
cat	1
mat	1
on	2
sat	2
sofa	1
the	4





## Example: Wordcount(2)

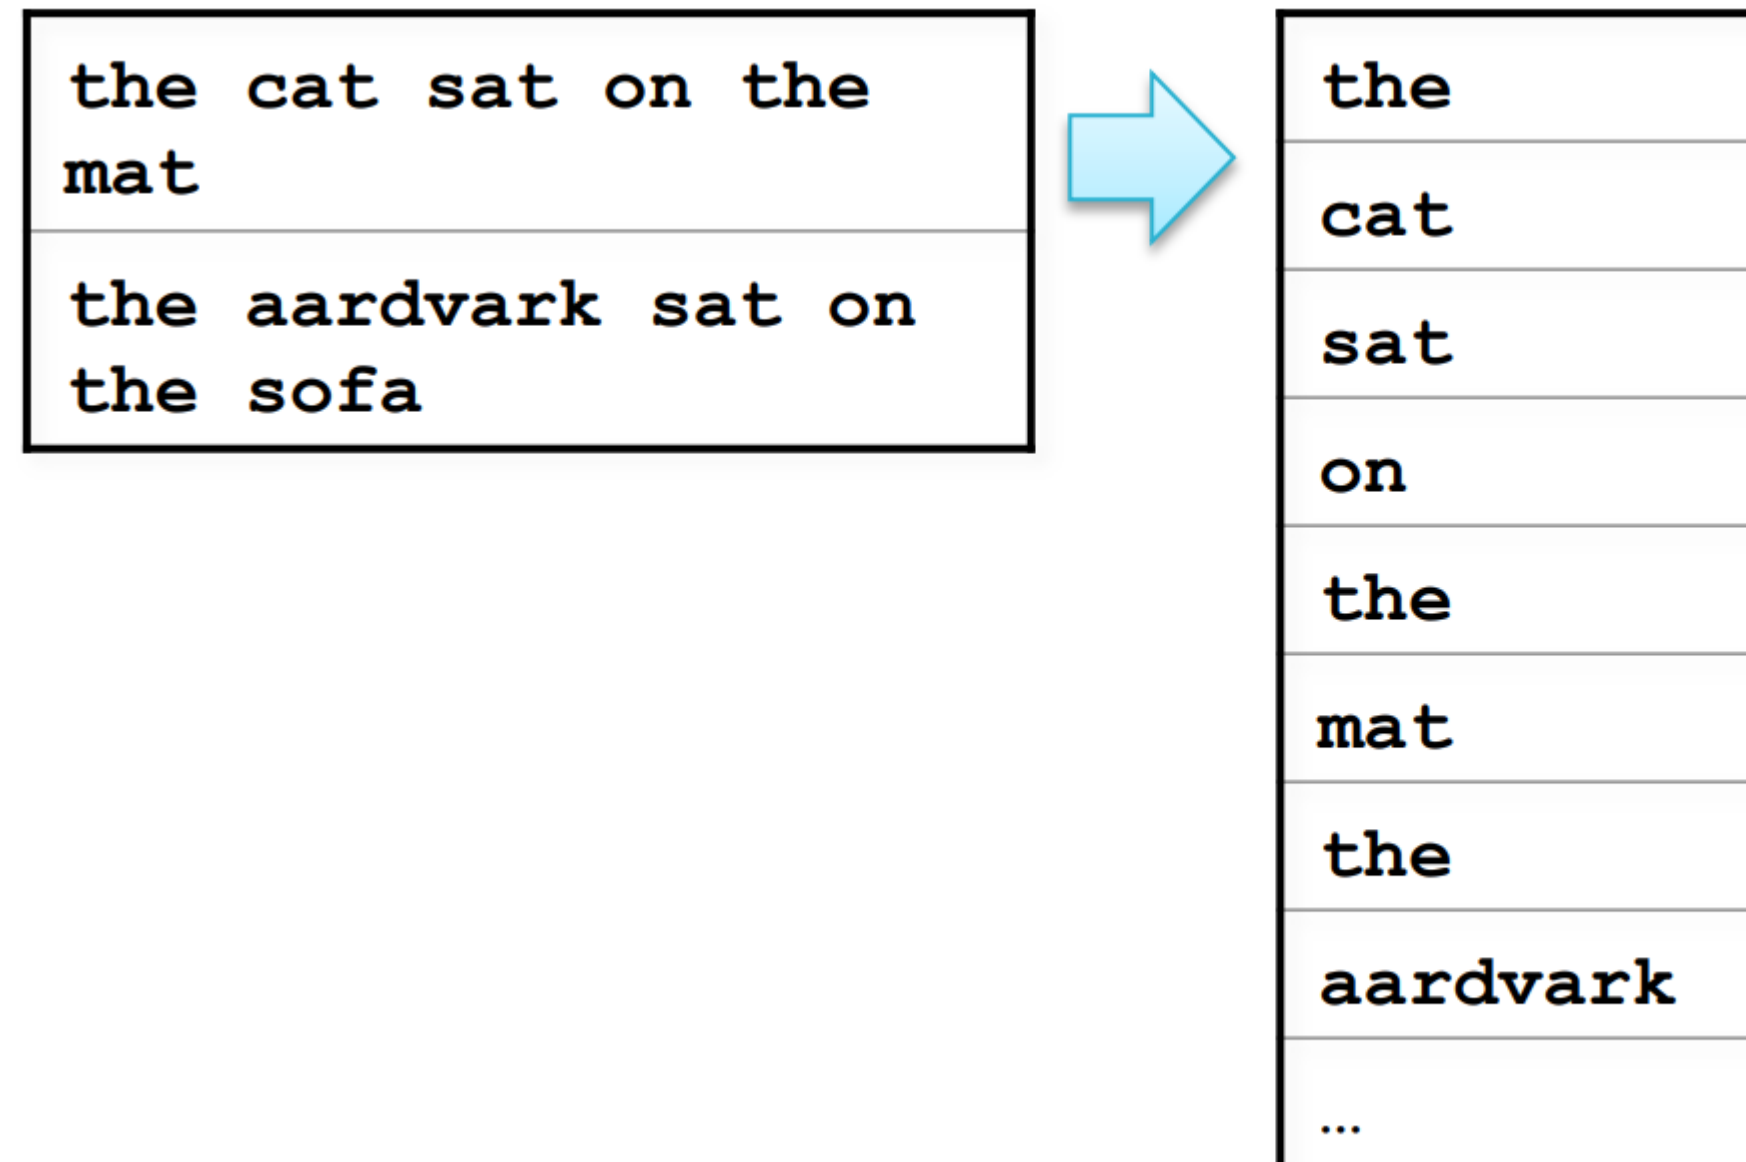
```
text_file = sc.textFile("hdfs://...")
```

```
the cat sat on the  
mat
```

```
the aardvark sat on  
the sofa
```

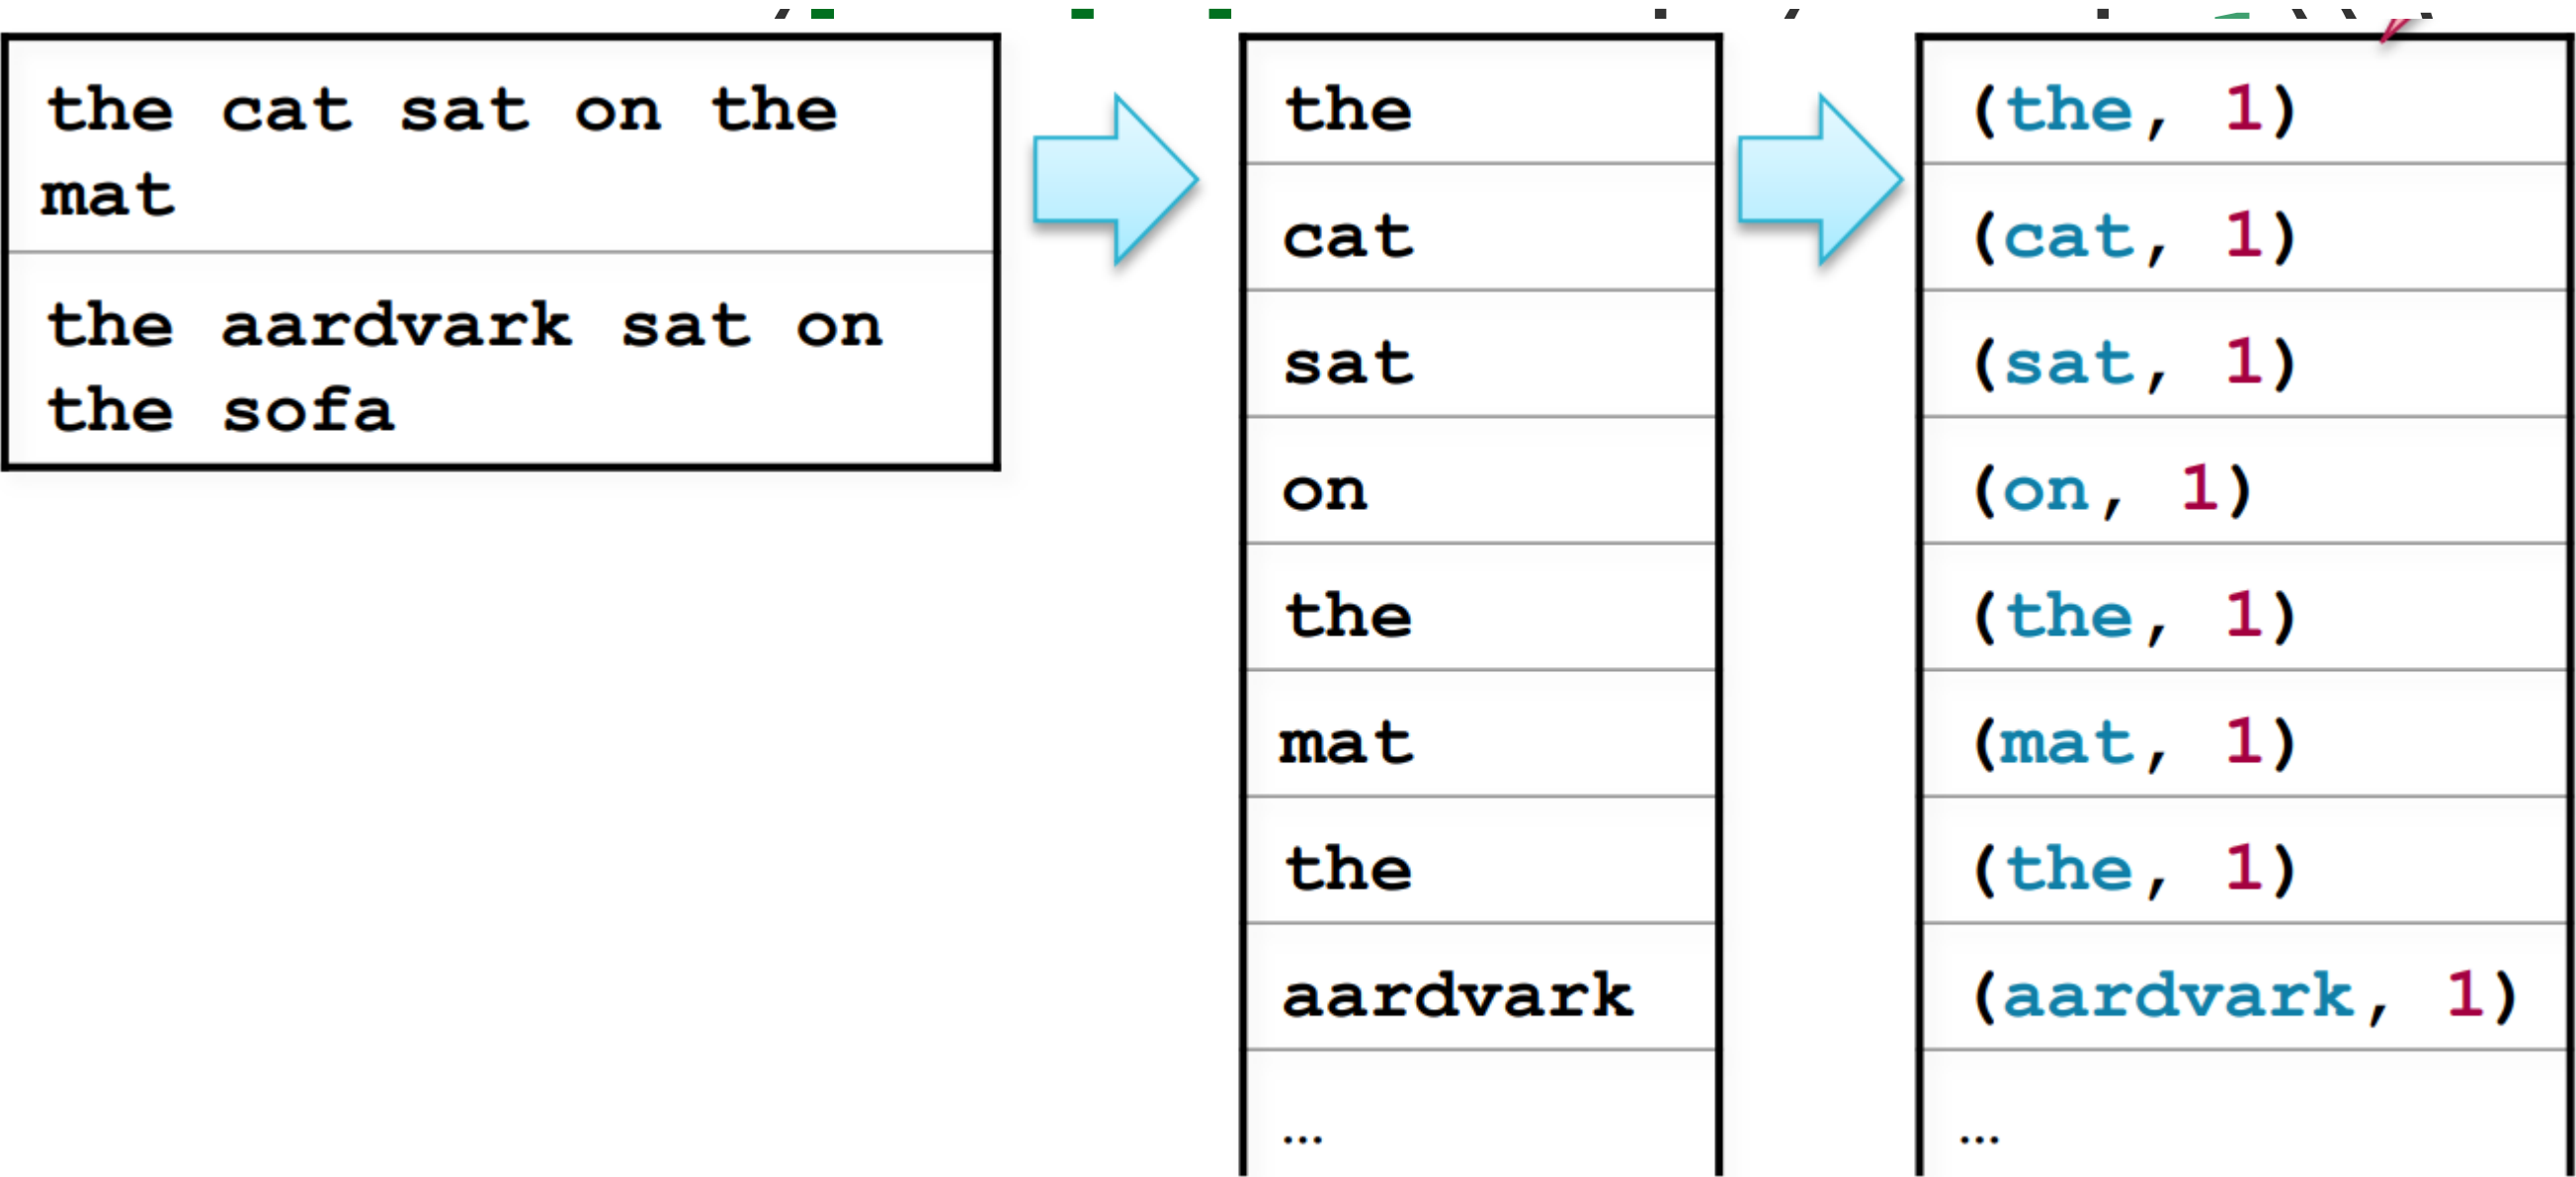
## Example: Wordcount(3)

```
text_file = sc.textFile("hdfs://...")  
counts = text_file.flatMap(lambda line: line.split(" "))
```



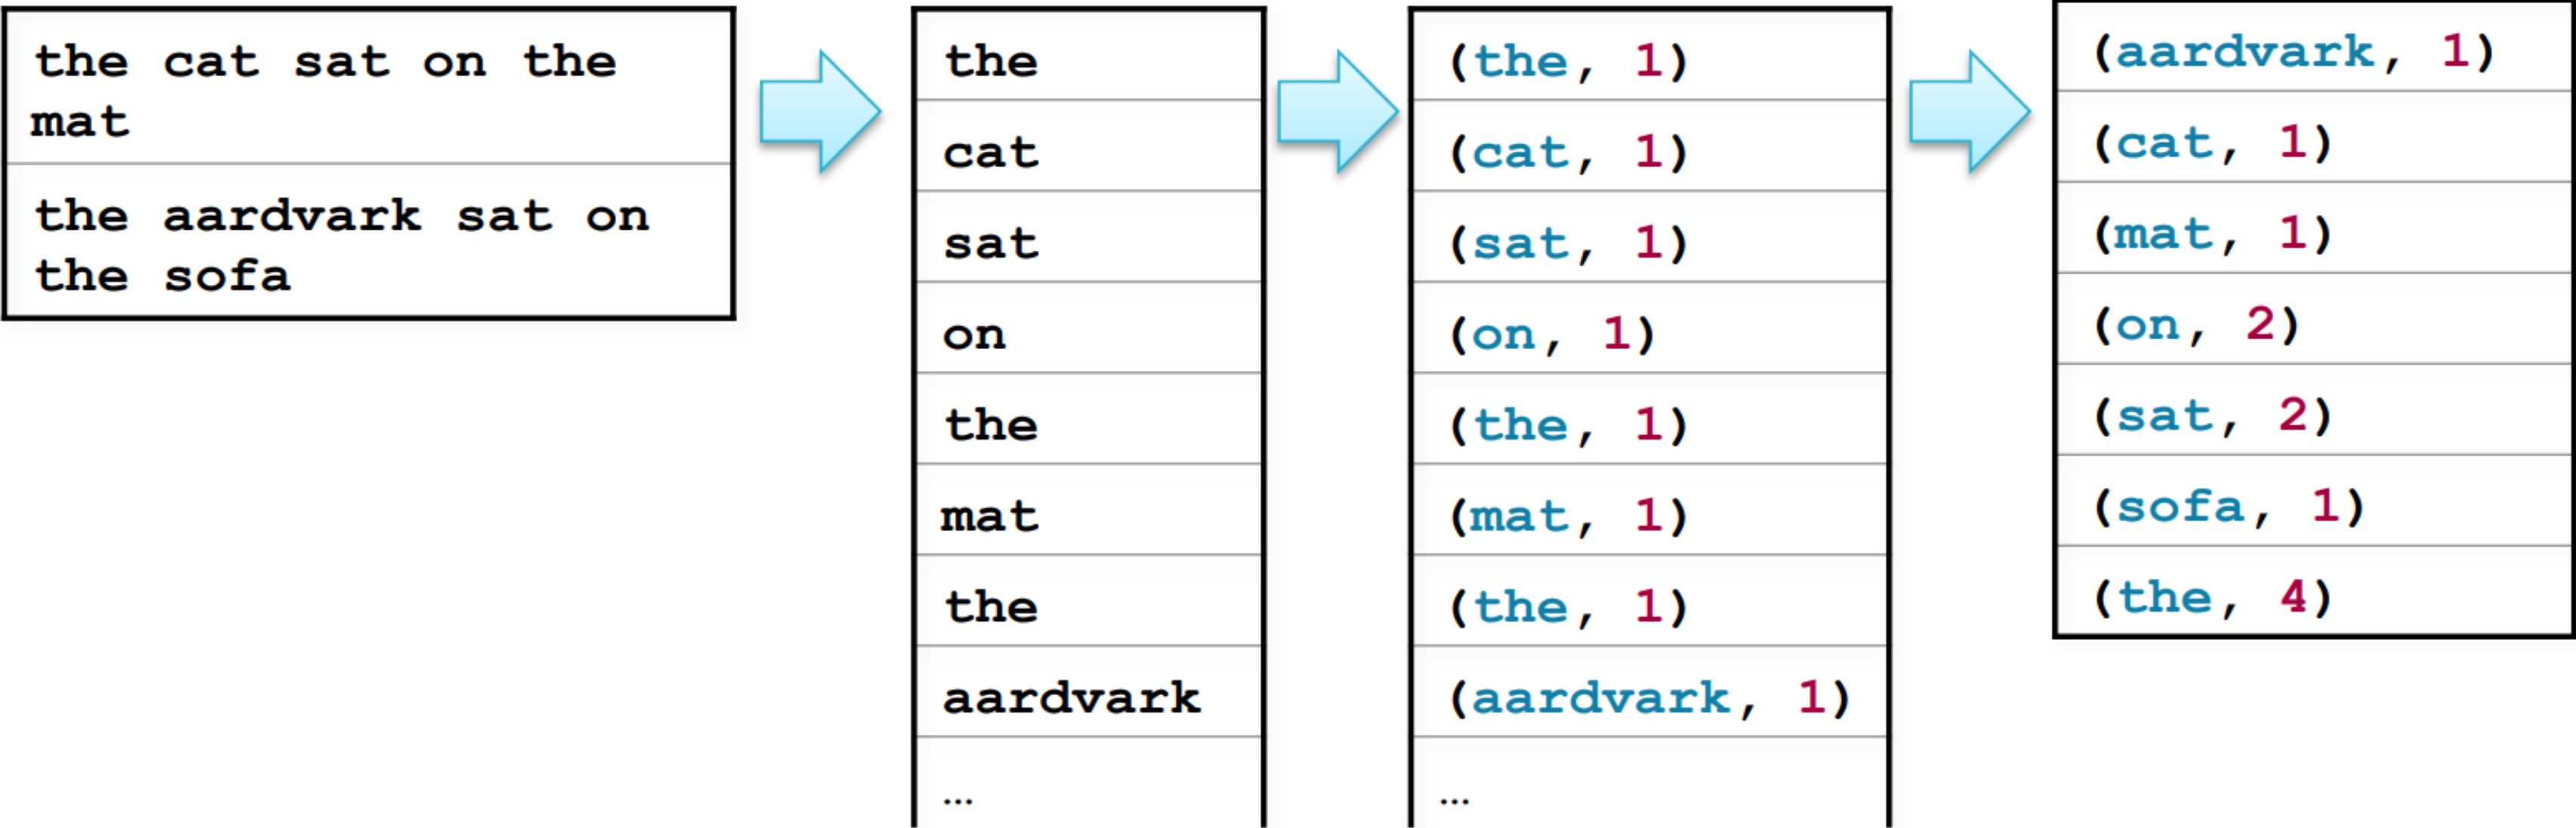
# Example: Wordcount(4)

```
text_file = sc.textFile("hdfs://...")
counts = text_file.flatMap(lambda line: line.split(" "))
\
```



# Example: Wordcount(5)

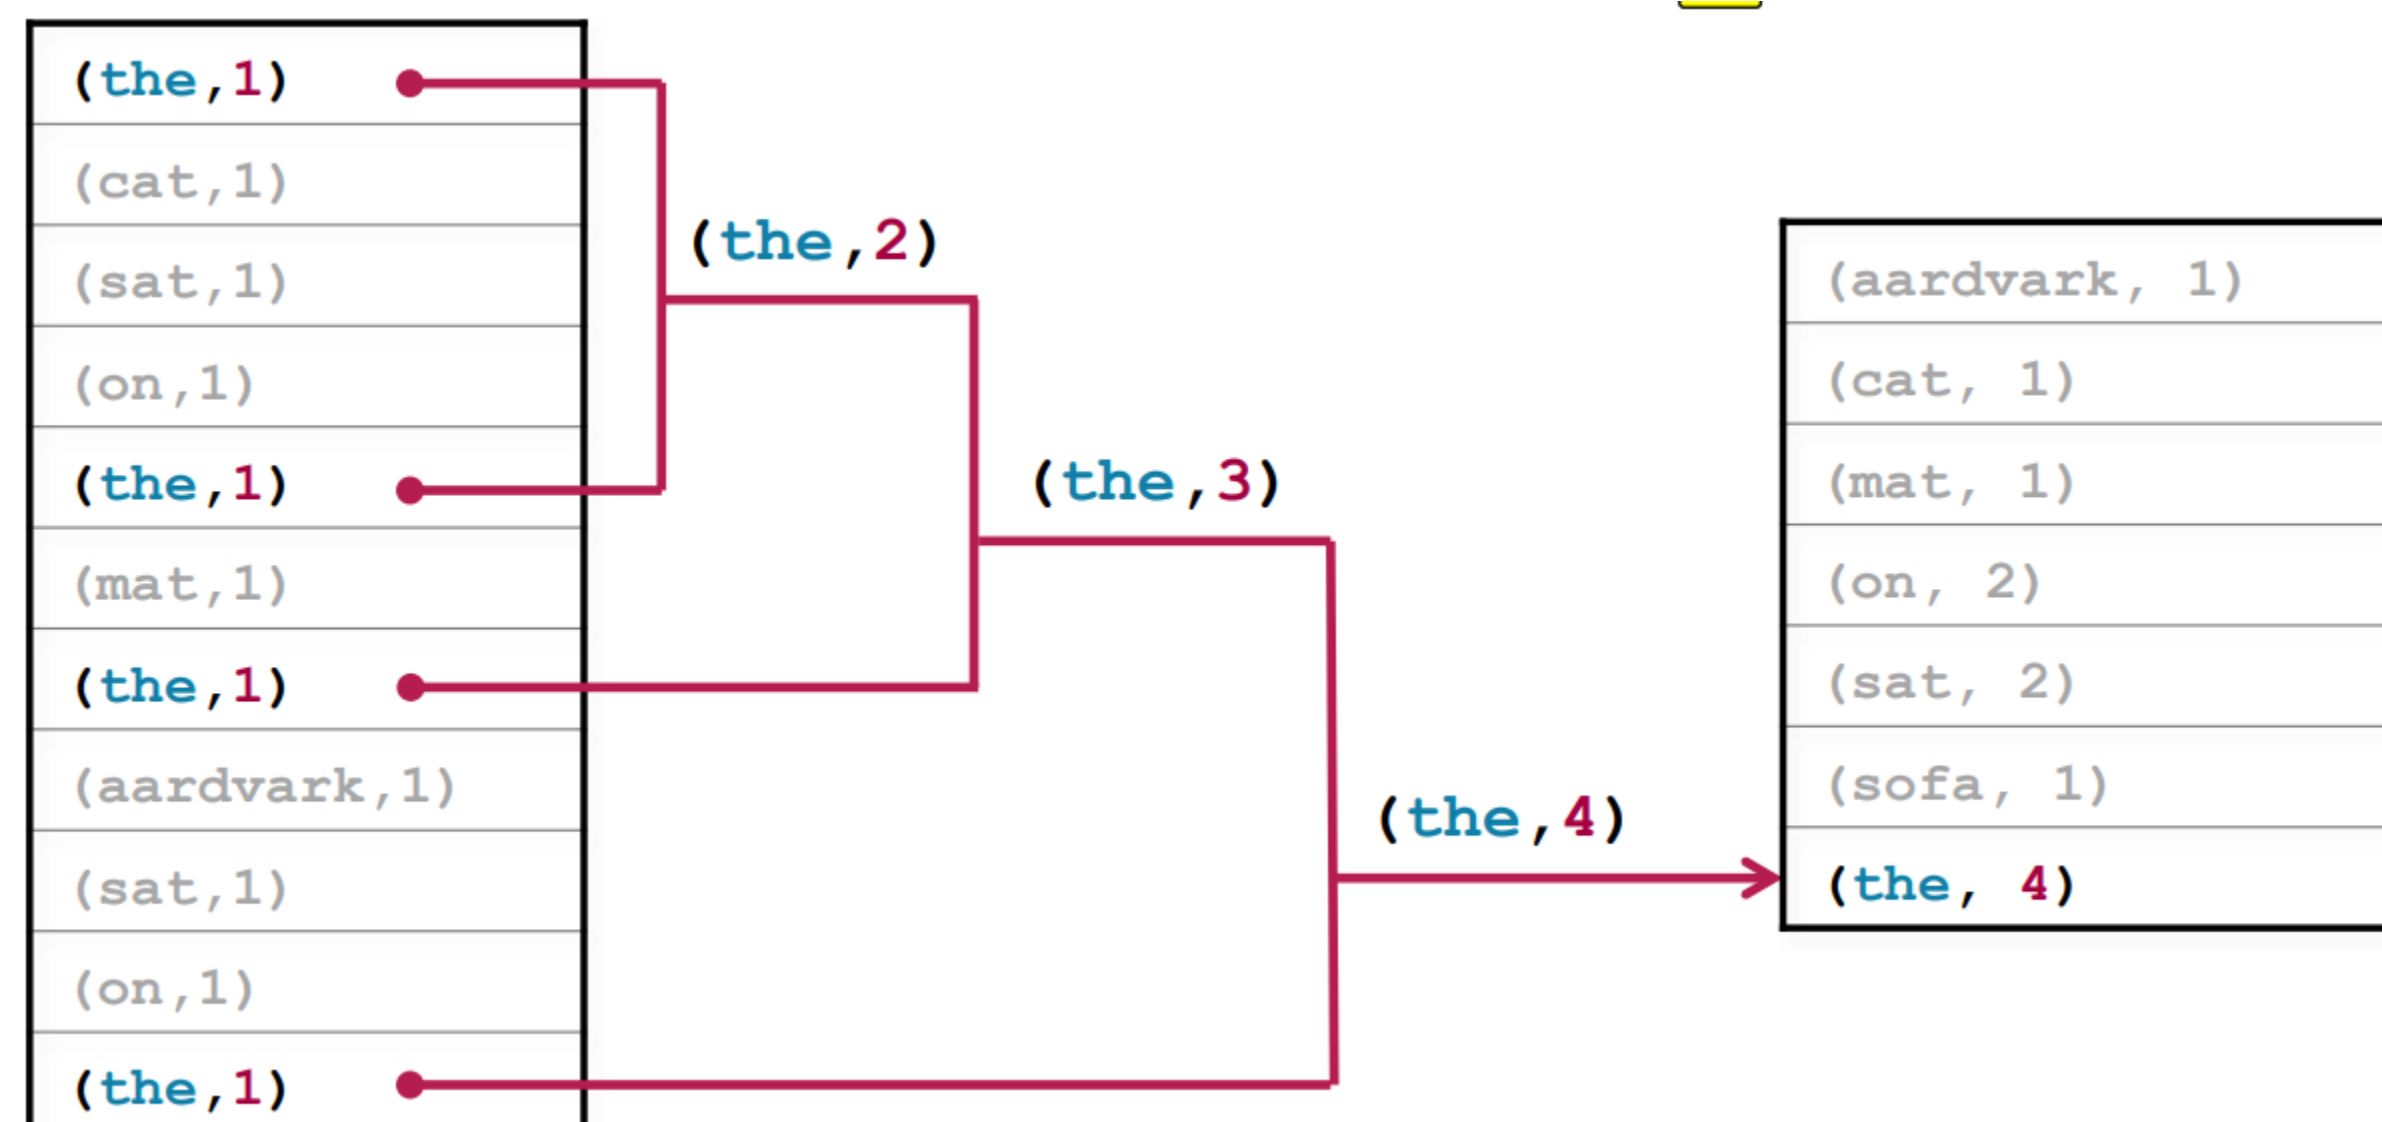
```
text_file = sc.textFile("hdfs://...")
counts = text_file.flatMap(lambda line: line.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)
counts.saveAsTextFile("hdfs://...")
```



## Example: Wordcount(5)

ReduceByKey functions must be

- Binary—combines values from two keys
- Commutative –  $x+y = y+x$
- Associative –  $(x+y)+z = x+(y+z)$



# Spark: use cases

- CERN: predict dataset popularity
- Verizon: feature extraction
- Uber: recommender system
- BBVA: tagging text in money transfers
- Toyota: power customer platforms
- Airbnb: predict demand